

# **When Front-Stage Innovation Overloads Back-Stage Workers: A Dynamic Model of Workforce, Capability, and Customer Trust in Omnichannel Retail**

(Authors' names are not included for peer review)

Authors are encouraged to submit new papers to INFORMS journals by means of a style file template, which includes the journal title. However, use of a template does not certify that the paper has been accepted for publication in the named journal. INFORMS journal templates are for the exclusive purpose of submitting to an INFORMS journal and are not intended to be a true representation of the article's final published form. Use of this template to distribute papers in print or online or to submit papers to another non-INFORM publication is prohibited.

**Abstract. Problem definition.** Digital investment in omnichannel retail generates back-stage demand on workers that standard dashboards cannot see. Front-stage automation amplifies operational complexity faster than back-stage capacity can absorb it, creating a structural measurement gap: technology returns and labour consequences appear on different ledgers, and the workforce that determines whether technology delivers its designed capability goes systematically underinvested.

**Methodology/results.** We formalise this gap as a parsimonious dynamic model and derive three analytically characterised results. (i) A computable digitalisation ceiling exists beyond which additional front-stage investment destroys customer lifetime value by overloading back-stage workers; front-stage demand metrics continue rising even as lifetime value collapses. (ii) The entire return to workforce investment flows through service capacity, invisible to dashboards that treat payroll as a cost line. (iii) The system is bistable: at our baseline calibration, the neglected firm faces a safe rollout window approximately 37% the size of its investing peer. Three empirical strategies provide directionally consistent evidence across different identification strategies. A statistically significant inverted-U between e-commerce share and customer satisfaction, and a post-shock decline in the American Customer Satisfaction Index following major omnichannel launches, support the digitalisation ceiling. County-level wage and employment data from the Bureau of Labor Statistics ( $N = 30,951$  county-year observations across 3,115 US counties, 2015–2024) support a labour-market analogue of the Good Jobs Buffer: counties with higher initial retail wages show significantly more resilient employment as e-commerce grows ( $p = 0.025$ ).

**Managerial implications.** The model supplies two computable numbers absent from standard dashboards: a digitalisation ceiling and a maximum safe rollout size. Both are workforce-contingent, making workforce investment a structural requirement of digital strategy rather than a social concession.

**Funding:** No external funding was received for this research.

**Key words:** omnichannel retail, workforce well-being, service capability, customer trust, system dynamics, people-centric operations

# 1. Introduction

Omnichannel grocery retailers have invested heavily in digital capabilities in recent years (Hübner et al. 2021, Bijmolt et al. 2021), enabling customers to build baskets across channels, reserve items for same-day collection, and track substitutions in real time. Yet as digital engagement has grown, retailers have encountered a structural asymmetry in how costs and returns are measured. Technology investment registers as successful in the front-stage domain where it is evaluated, while the back-stage workforce that must fulfil the digitally generated demand absorbs consequences that appear nowhere in the technology business case (Corbett 2024).

The asymmetry is structural, not incidental. Front-stage digital capabilities multiply handoff complexity far faster than order volume grows (Saghiri et al. 2017), placing distinctive demands on back-stage workers. Workforce investment generates returns that are large and invisible to standard metrics (Fisher et al. 2021, Kesavan et al. 2022, Ton 2014), yet the causal chain runs through mechanisms that dashboards do not monitor.

What is missing is a formal account of the mechanism. The empirical literature establishes that workforce investment pays off; it does not characterise *when* front-stage investment begins to destroy value, *how fast* a firm can scale digitalisation before the workforce is overwhelmed, or *what dynamic relationship* connects technology, well-being, and trust. These questions require a model of co-evolving state variables responding to a controllable investment sequence (see §2).

This paper formalises the economic structure of this problem as a parsimonious dynamic model. Technology and workforce are not independent inputs to service capability but strategic complements (Milgrom and Roberts 1990, Bainbridge 1983); in the spirit of Kremer (1993), where the weakest link constrains system output, the capacity multiplier formalises the dependence of technology delivery on workforce engagement. The model nests the quality-erosion framework of Oliva and Sterman (2001) as a special case (EC §EC.2). The central modelling innovation is the *workforce capacity multiplier*: effective service capability is the product of the technology stock and a workforce engagement index. When front-stage demand grows faster than back-stage workforce capacity, a service gap opens that depletes workforce well-being, which reduces

effective capability, which widens the gap further: a reinforcing deterioration that single-domain models cannot generate.

Three formally characterised results follow. A *Front-Stage Trap* (Theorem 1): a computable digitalisation threshold exists beyond which additional front-stage investment reduces customer lifetime value by overloading back-stage workers. A *Good Jobs Buffer* (Theorem 2): the return to workforce investment flows through a capacity-level channel that is invisible to dashboards measuring only operational sensing. A *Burnout Cliff* (Theorem 3): the system exhibits bistability, and at baseline calibration, a firm that neglects its workforce faces a safe deployment window approximately 37% the size of one that invests in its people. A fourth result, Proposition 4, establishes that the marginal value of workforce investment is increasing in digitalisation level; Corollary 5 (a consequence of Theorems 1 and 2, extended computationally) shows that the all-technology budget allocation is dominated at and near the Trap threshold, and computationally throughout the tested range. Workforce investment is therefore not a social concession but a structural requirement of digital strategy.

## 2. Related Literature

Three bodies of research motivate this work, each contributing part of the picture but none assembling it. We review each in turn.

*Service quality erosion.* The foundational insight that service quality can erode endogenously under demand pressure originates with Oliva and Sterman (2001), whose system-dynamics model (Sterman 2000) produces death spirals in which workers cut corners under demand pressure. Their finding that workers are roughly twice as willing to cut corners as to work overtime grounds the overload-asymmetry parameter  $\sigma$  in the present model. Subsequent work extends the logic to capability erosion (Rahmandad and Repenning 2016, Repenning and Sterman 2001), e-commerce growth (Oliva et al. 2003), and regulatory compliance (Martinez-Moyano et al. 2014). KC and Terwiesch (2009) confirm empirically, in a hospital context, that sustained high workload degrades service outcomes—a mechanism this model formalises for retail. Corbett (2024) calls for formal dynamic models of workforce well-being in operations.

**Gap 1:** no model in this tradition treats workforce well-being as a positive capability multiplier. **Gap 2:** no Oliva–Sterman extension adds customer trust, a front-/back-stage split, or digital scaling.

*Workforce investment and retail performance.* A substantial empirical literature documents the workforce-performance link. The Good Jobs Strategy (Ton 2014) identifies four operational choices that drive outperformance; causal evidence shows that responsible scheduling raises productivity (Kesavan et al. 2022) and that profit-maximising staffing exceeds prevailing practice (Fisher et al. 2021, Dubé et al. 2020). The conceptual mechanism is the *service profit chain* (Heskett et al. 1994), which has not been formalised as a dynamic model with computable boundary conditions. The Job Demands-Resources (JD-R) model (Bakker and Demerouti 2007, Demerouti et al. 2001) provides the micro-foundation: overload depletes well-being faster than equivalent surplus restores it, and productivity scales with engagement (Mas and Moretti 2009).

Prospect theory (Kahneman and Tversky 1979, Tversky and Kahneman 1992) provides a deeper behavioural foundation: losses loom larger than equivalent gains by  $\lambda \approx 2.25$ , and both domains exhibit diminishing sensitivity. Applied to the workforce, overload-induced depletion is experienced more intensely than equivalent recovery—a prediction remarkably consistent with the 2:1 asymmetry in Oliva and Sterman (2001) and the JD-R model. Loss aversion thus independently grounds the parameter  $\sigma > 1$  central to Theorems 1 and 3.

**Gap 3:** no dynamic model formalises the Good Jobs Strategy or the service profit chain with computable boundary conditions.

*Omnichannel capability.* Digitalisation creates considerable challenges for retail logistics, with fulfilment complexity scaling super-linearly as channel combinations increase (Hübner et al. 2021, Saghir et al. 2017, Bell et al. 2018, Acimovic and Graves 2015). Automation shifts the operator's role toward exception handling (Bainbridge 1983), making the quality of residual human tasks more consequential for overall output (Gans and Goldfarb 2026). Technology–workforce strategic complementarity (Milgrom and Roberts 1990, Aral et al. 2012) motivates the model's central structural assumption that the marginal value of workforce engagement rises with the technology base (§3.1). **Gap 4:** no model formalises front-to-back cross-domain overload. **Gap 5:** the digitalisation-workforce cross-partial has not been derived for services.

Table 1 summarises the five gaps and the model's response to each.

**Table 1 Literature gaps and model contributions.**

| Gap                                                                        | Closest prior work                                                  | This paper                                                                          |
|----------------------------------------------------------------------------|---------------------------------------------------------------------|-------------------------------------------------------------------------------------|
| 1. No model treats $W$ as a positive capability multiplier                 | Oliva and Sterman (2001): $W$ as degradation only                   | Effective capability as the product of technology and workforce engagement          |
| 2. No Oliva–Sterman extension adds trust, domain split, or digital scaling | Oliva and Sterman (2001); Martinez-Moyano et al. (2014)             | Customer trust as a co-evolving state; front/back domain split; digital scaling     |
| 3. Good Jobs Strategy / service profit chain (SPC) unformalised            | Ton (2014); Heskett et al. (1994): empirical only                   | Dynamic formalisation with analytical equilibria and boundary conditions            |
| 4. No cross-domain front-to-back overload model                            | Saghiri et al. (2017): conceptual only                              | Super-linear back-stage demand as a function of digital depth and customer trust    |
| 5. Digitalisation-workforce cross-partial not derived for services         | No direct service prior; see also Gans and Goldfarb (2026) (O-ring) | Formally derived: the return to workforce investment rises with the technology base |

### 3. Model

The model has one moving part that prior frameworks lack: effective capability depends on the workforce. The front-stage digitalisation level  $F > 0$  is an exogenous managerial control that sets digital feature complexity and thereby determines back-stage demand. The model traces the consequences of choosing  $F$  for three co-evolving state variables and, through them, for customer lifetime value ( $CLV =$  a standard increasing function of trust; calibrated in §5, but all theorems require only  $\partial CLV / \partial T^* > 0$ ). Each of the eight free parameters has a single role with a clean kill condition (Table 2).

#### 3.1. Stocks and Effective Capability

Three state variables evolve jointly. The first,  $K(t) > 0$ , is the back-stage technology stock: installed infrastructure, warehouse management systems, and process standardisation. The second,  $W(t) \in (0, 1)$ , indexes workforce well-being—the engagement, energy, and cognitive capacity of pickers, drivers, and click-and-collect associates. The third,  $T(t) \in (0, 1)$ , measures customer trust: the aggregate confidence that the retailer will deliver reliably. All three are dimensionless and normalised to their natural operating ranges. Throughout, starred variables ( $K^*, W^*, T^*$ ) denote steady-state equilibrium values;  $F_0$  is the pre-shock digitalisation level;  $\Delta F = F - F_0$  is the upgrade step.

The central structural innovation is the *effective capability*. A warehouse management system or substitution-handling process can only deliver its designed performance if the workforce operating it has

the capacity to do so. Effective capability is therefore not the technology stock alone but its product with a workforce engagement multiplier:

$$K_{\text{eff}}(t) = K(t) \cdot h(W(t)),$$

$$h(W) = (1 - p) + pW, \quad p \in (0, 1). \quad (1)$$

Effective capability is the product of the technology stock and a workforce capacity multiplier  $h(W)$ , ranging from  $h(0) = 1 - p$  (zero well-being: only the automation-driven fraction is realised) to  $h(1) = 1$  (full design capacity). The cross-partial

$$\frac{\partial^2 K_{\text{eff}}}{\partial K \partial W} = p > 0 \quad (2)$$

is strictly positive for all  $p > 0$ , establishing technology and workforce as complements in the Milgrom and Roberts (1990) sense.

The parameter  $p$  governs complementarity:  $p = 0$  decouples workforce from capability;  $p = 1$  makes the operation fully labour-intensive. The baseline  $p = 0.5$  is an illustrative moderate starting point; estimation is discussed in §7.2. All three theorems hold qualitatively for any  $p \in (0, 1)$ .

### 3.2. Demand and the Service Gap

Customer demand on back-stage operations is

$$D(T, F) = (1 + aT) \cdot F^b, \quad (3)$$

where  $a > 0$  is the trust-demand coupling and  $b > 1$  is the front-stage amplification exponent. The trust-demand coupling reflects the established finding that customer trust drives spending consolidation and order frequency (Sirdeshmukh et al. 2002, McKnight et al. 2002). The super-linearity  $b > 1$  captures the well-documented increase in back-stage fulfilment complexity as front-stage channel breadth grows: as shown in §2, Bell et al. (2018) and Acimovic and Graves (2015) document patterns consistent with fulfilment error rates and demand-capacity mismatches rising faster than linearly with channel combinations. We set  $b = 1.3$  as an illustrative, structurally conservative baseline motivated by this literature; the Front-Stage Trap requires only  $b > 1$ , and all three theorems hold qualitatively across  $b \in [1.1, 2.0]$  (EC Table EC.3).

The *service gap* measures the shortfall between demand and effective capability:

$$G(t) = D(T(t), F) - K(t) \cdot h(W(t)). \quad (4)$$

When  $G > 0$  (overload), customers experience failures. When  $G < 0$  (surplus), service runs reliably. We write  $[G]^+ = \max(G, 0)$  and  $[-G]^+ = \max(-G, 0)$  for the positive and negative parts.

### 3.3. Dynamics

Three ordinary differential equations translate the overload and surplus mechanisms into rate equations. Under overload, workers cut corners, trust erodes, and capability may spiral downward; under surplus, capacity allows recovery. All rate parameters carry units of  $\text{months}^{-1}$ .

*Capability.* The technology stock evolves as

$$\frac{dK}{dt} = r [G]^+ W - \delta K + u, \quad (5)$$

where  $r > 0$  is the sensing-adaptation rate,  $\delta > 0$  is the depreciation rate, and  $u \geq 0$  is direct back-stage investment. The sensing term  $r [G]^+ W$  captures organisational learning from service failures (Tucker 2007): the product structure requires both a gap to detect ( $[G]^+ > 0$ ) and workforce capacity to act on it ( $W > 0$ ), reflecting the finding that capability accumulation scales with the workforce knowledge stock (Argote and Epple 1990). Under surplus ( $[G]^+ = 0$ ), the sensing term vanishes and the firm accumulates no capability from successful operations; this asymmetry does not qualitatively alter the results because all three theorems are driven by overload-regime feedback.

*Workforce well-being.* Well-being rises under surplus capacity and direct investment; it falls through natural decay and overload-amplified depletion:

$$\frac{dW}{dt} = ([-G]^+ + f)(1 - W) - (\mu + \sigma [G]^+) W. \quad (6)$$

Two grouped terms ensure  $W(t) \in [0, 1]$ . The first drives  $W$  upward (vanishing at  $W = 1$ ): capability surplus and direct workforce investment raise well-being with diminishing returns, consistent with the JD-R model (Bakker and Demerouti 2007). The second drives  $W$  downward (vanishing at  $W = 0$ ): natural decay ( $\mu$ ) and overload erosion ( $\sigma [G]^+$ ) deplete well-being, with  $\sigma > 1$  reflecting the asymmetry documented by Oliva and Sterman (2001) and independently predicted by the loss-aversion ratio  $\lambda \approx 2.25$  from prospect theory (Tversky and Kahneman 1992). All three theorems hold for any  $\sigma > 1$  (EC Table EC.3).



*Customer trust.* Trust builds through reliable service and erodes faster when reliability fails (Sirdeshmukh et al. 2002, Mittal et al. 1998):

$$\frac{dT}{dt} = \beta [-G]^+ (1 - T) - \eta \beta [G]^+ T - \lambda T, \quad (7)$$

where  $\beta > 0$  is the trust-update rate,  $\eta > 1$  is the erosion asymmetry, and  $\lambda > 0$  is natural trust decay. The  $\lambda T$  term captures gradual trust erosion absent service experience—customer-base turnover, competitive switching, and memory decay (McKnight et al. 2002)—with a calibrated half-life of approximately 28 months, conservative relative to omnichannel grocery churn rates; all qualitative results hold for any  $\lambda > 0$ . We set  $\eta = 2.0$ , below the approximately 2.5:1 ratio reported by Sirdeshmukh et al. (2002).

With all three ODEs specified, we note a technical property. The system (5)–(7) is piecewise-smooth; existence and uniqueness of solutions follow from Filippov's (1988) framework (details and smoothing robustness in EC §EC.1.1 and §EC.1.8).

### 3.4. State-Space Invariance and Empirical Mapping

The state space  $K > 0$ ,  $W \in [0, 1]$ ,  $T \in [0, 1]$  is forward-invariant for all  $t \geq 0$ ; the verification via directional derivatives at each boundary is given in EC §EC.1.1.


Parameter calibration details are in EC §EC.4.

### 3.5. Feedback-Loop Architecture

Figure 1 shows the causal-loop architecture of the three ODEs, using the conventions of Sterman (2000). The diagram contains four feedback loops, each traceable from the state-variable arrows.

**Loop R1** ( $T \rightarrow D \rightarrow G \rightarrow T$ ) has regime-dependent polarity. Under surplus, it is balancing: trust-driven demand growth is absorbed by excess capacity. Under overload, it turns reinforcing: trust loss reduces demand, but capacity collapses faster, widening the gap. **Loop R2** ( $G \rightarrow W \rightarrow K_{\text{eff}} \rightarrow G$ ) is reinforcing: a wider service gap depletes the workforce, which reduces effective capability, which widens the gap further. This is the deterioration loop behind the Front-Stage Trap. **Loop R3** ( $W \rightarrow K_{\text{eff}} \rightarrow G \rightarrow W$ ) is also reinforcing: it is the escalation channel behind the Burnout Cliff. **Loop B1** ( $G \rightarrow K \rightarrow K_{\text{eff}} \rightarrow G$ ) is balancing: sensing investment closes the gap, but only when the workforce has the capacity to act (see the  $r[G]^+W$

**Figure 1 Causal-loop architecture. Solid boxes are state variables; dashed boxes are exogenous controls. Arrow signs indicate direction of effect; loop polarity is shown in parentheses. The  $(\pm)$  notation for R1 indicates regime-dependent polarity: balancing when  $G < 0$  (surplus), reinforcing when  $G > 0$  (overload). R2 and R3, created by the capacity multiplier  $h(W)$ , are novel to this model.**



term in Eq. 5). Loops R2 and R3 are novel to this model. They have no counterpart in Oliva and Sterman's (2001) two-loop structure (which contains only B1 and a variant of R1), and they are the source of the regime-dependent instability that drives all three theorems.

### 3.6. Parameters and Special Cases

Table 2 summarises all parameters. Two control variables complete the model:  $u \geq 0$  (capability investment) and  $f \geq 0$  (workforce investment). The model recovers the core dynamics of Oliva and Sterman (2001) when  $p = 0$  and  $a = 0$  (see EC §EC.2 for the structural mapping), and reduces to a Solow-type capital-accumulation equation when  $p = 0$ ,  $a = 0$ , and  $r = 0$ .

## 4. Analytical Results

Four results follow from the model, each identifying a phenomenon that single-domain frameworks cannot produce. Theorems 1–3 characterise three such phenomena; Proposition 4 establishes digitalisation-workforce complementarity. Full proofs are in the Electronic Companion (§EC.1); here we provide proof sketches.

### 4.1. Equilibrium Structure

The model generates two distinct long-run outcomes. In the first, which we call the *surplus equilibrium*, effective capability exceeds demand. The workforce is healthy, trust builds, and the firm settles into a self-reinforcing high-performance state. In the second, the *low-trust attractor*, demand has overwhelmed capability, the workforce is depleted, and customers have withdrawn trust. Exit from this attractor requires structural

**Table 2** Model parameters, baseline values, and kill conditions.

| # | Parameter                       | Baseline | Timescale                    | Kill                  | Mechanism killed     | Empirical anchor                                            |
|---|---------------------------------|----------|------------------------------|-----------------------|----------------------|-------------------------------------------------------------|
| 1 | $b$ (demand super-linearity)    | 1.3      | –                            | $b=1$                 | Front-Stage Trap     | Bell et al. (2018); Acimovic & Graves (2015)                |
| 2 | $a$ (trust-demand coupling)     | 0.4      | –                            | $a=0$                 | Demand feedback      | McKnight et al. (2002); Sirdeshmukh et al. (2002)           |
| 3 | $p$ (workforce complementarity) | 0.5      | –                            | $p=0$                 | Capacity multiplier  | Illustrative                                                |
| 4 | $\sigma$ (overload asymmetry)   | 2.0      | –                            | $\sigma=1$            | Tipping point        | Oliva and Sterman (2001); PT $\lambda \approx 2.25^\dagger$ |
| 5 | $\eta$ (erosion asymmetry)      | 2.0      | –                            | $\eta=1$              | Trust asymmetry      | Sirdeshmukh et al. (2002)                                   |
| 6 | $r$ (sensing-adaptation)        | 0.4      | $\tau_r \approx 1.7$ mo      | $r=0$                 | Self-correction dies | Tucker (2007); Argote & Epple (1990)                        |
| 7 | $\delta$ (depreciation)         | 0.03     | $\tau_\delta \approx 23$ mo  | –                     | –                    | Illustrative (EC §EC.4)                                     |
| 8 | $\beta$ (trust-update rate)     | 0.12     | $\tau_\beta \approx 5.8$ mo  | $\beta \rightarrow 0$ | Trust frozen         | Illustrative (EC §EC.4)                                     |
| + | $\mu$ (natural $W$ decay)       | 0.05     | $\tau_\mu \approx 14$ mo     | $\mu=0$               | $W$ exponential only | Demerouti et al. (2001); Dubé et al. (2020)                 |
| + | $\lambda$ (natural $T$ decay)   | 0.025    | $\tau_\lambda \approx 28$ mo | $\lambda=0$           | $T$ exponential only | Illustrative (EC §EC.4)                                     |

All time units are months. Timescales are  $\ln(2)/\theta$  (half-life). Eight free parameters (rows 1–8);  $\mu$  and  $\lambda$  govern natural decay. “Illustrative” parameters are calibrated within empirically informed ranges (EC §EC.4).  $^\dagger$  Magnitude comparator from service industry evidence; see §3.3. Direct estimation of  $b$  and  $p$  from retailer panel data is the primary empirical extension (§7.2).

intervention: because  $r[G]^+W \rightarrow 0$  as  $W \rightarrow 0$ , the firm cannot learn its way out, and recovery requires direct investment in both  $u$  and  $f$ .

The formal conditions for each state are as follows. The surplus equilibrium exists for all  $F < F_{\text{crit}}$  and is unique under Condition C2 (EC Lemma 2, §EC.1.2). When effective capability exceeds demand, define the operational surplus  $S^* \equiv K^*h(W^*) - D(T^*, F) \geq 0$  as the excess of effective capability over demand at equilibrium. Equilibrium trust is an increasing function of  $S^*$ , saturating toward one as surplus grows:  $T^* = \beta S^* / (\beta S^* + \lambda)$ . Equilibrium well-being is similarly increasing in surplus and investment:  $W^* = (S^* + f) / (S^* + f + \mu)$ . The capability stock settles at  $K^* = u/\delta$ , determined entirely by the investment-depreciation balance. Full algebraic derivations are in EC §EC.1.2. Both  $T^*$  and  $W^*$  are increasing in  $S^*$ , so any improvement in effective capability propagates into higher trust and well-being at the equilibrium, creating the reinforcing dynamics that underpin the Good Jobs Buffer.

The low-trust attractor is reached whenever demand persistently overwhelms capability. With  $T^* \approx 0$ , demand collapses to its digitalisation-driven floor and the workforce settles at a low but strictly positive well-being level, held above zero by the workforce investment rate  $f$ . Effective capability at this attractor falls short of demand, maintaining a persistent residual gap. The attractor exists whenever  $F^b > (u/\delta)h(W_{\text{low}}^*)$  and is locally stable (EC Remark 3).

**Maintained conditions.** The formal results hold under four regularity conditions (C1)–(C4), verified analytically or numerically across the parameter space (EC Remark 5): surplus equilibrium existence, uniqueness, regularity, and low-trust attractor existence (EC §EC.1.2). The coexistence of surplus equilibrium and low-trust attractor constitutes bistability; the transition at  $F_{\text{crit}}$  where these two states coalesce is classified as a saddle-node (fold) bifurcation (EC Remark 6).

**REMARK 1 (ANALYTICAL TRANSPARENCY).** The table below classifies each result by evidentiary basis: Tier A (analytic for all valid parameters), Tier B (analytic under C1–C4), or Tier C (computational verification).

| Result                                                 | Tier | Basis                                                                                           |
|--------------------------------------------------------|------|-------------------------------------------------------------------------------------------------|
| Thm 1(i): equilibrium existence/nonexistence           | A    | Analytic (IVT + $\Phi(0)$ sign)                                                                 |
| Thm 1(ii): trajectories bounded                        | A    | Analytic (Gronwall; Lemma 3)                                                                    |
| Thm 1(ii): wide-basin convergence                      | C    | Numerical grid, $K_0 = K^*$ slice                                                               |
| Lemma 2 (uniqueness of $S^*$ )                         | B    | Analytic under C2-SC ( $\beta(f+\mu) < \lambda$ ; holds throughout parameter space; EC §EC.1.2) |
| Thm 2: $\partial \text{CLV}^* / \partial f > 0$ (sign) | B    | IFT chain; GE feedback correction sign analytic                                                 |
| Thm 2: GE feedback magnitude ( $\approx 0.09$ )        | C    | Confirmed across 100 configurations                                                             |
| Thm 3(i): $\Delta F_{\text{collapse}}$ threshold       | A    | Direct from Thm 1                                                                               |
| Thm 3(i): convergence from pre-shock IC                | C    | Numerical; basin-dependent                                                                      |
| Prop 4 cross-partial                                   | A    | Analytic ( $K_{\text{eff}} = Kh(W)$ )                                                           |
| Prop 4 $F_{\text{max}}(W)$ boundary formula            | A    | Exact at $S^* \rightarrow 0$                                                                    |
| Prop 4 $F_{\text{max}}(W)$ interior upper bound        | B    | Analytic sign; $\approx 8\%$ correction numeric                                                 |
| Corollary 5, Case 1                                    | B    | Thm 1 + Thm 2 directly                                                                          |
| Corollary 5, Case 2 near $B_{\text{min}}$              | B    | Continuity from Case 1 (IFT; EC §EC.1.7)                                                        |
| Corollary 5, Case 2 interior                           | C    | Computational; $c_F/c_f \in [0.5, 5]$                                                           |

## 4.2. Theorem 1: The Front-Stage Trap

**THEOREM 1 (Front-Stage Trap).** Suppose  $b > 1$  and  $p > 0$ . Define the digitalisation threshold

$$F_{\text{crit}} = \left( \frac{u}{\delta} \cdot \left[ (1-p) + p \frac{f}{f+\mu} \right] \right)^{1/b}, \quad (8)$$

which depends only on exogenous parameters  $(u, \delta, p, f, \mu, b)$  and equals  $K^*h(f/(f + \mu))^{1/b}$ , the capacity multiplier evaluated at the boundary workforce level. Then:

(i) For  $F < F_{\text{crit}}$ , the surplus equilibrium (unique under C2; Lemma 2, EC §EC.1.2) exists with  $T^* > 0$  and  $\text{CLV}^* > \text{CLV}_{\text{low}}^1$  (the CLV at the low-trust attractor).

(ii) For  $F \geq F_{\text{crit}}$ , no surplus equilibrium exists. All forward trajectories remain in a bounded compact set (Lemma 3, EC §EC.1.2), and the low-trust attractor is locally stable (EC Remark 3). Wide-basin convergence to this attractor is confirmed numerically (Tier C; Remark 1).


The threshold  $F_{\text{crit}}$  is decreasing in  $p$  and increasing in  $f$ : the Trap binds earlier when workforce-technology coupling is strong (because a more coupled system depends more heavily on the boundary well-being level  $f/(f + \mu) < 1$ , reducing effective capacity at the threshold), and is relaxed by workforce investment.

*Proof sketch.* The surplus equilibrium requires  $\Phi(0) > 0$  (where  $\Phi(S) \equiv K^*h(W^*(S)) - D(T^*(S), F) - S$  is the surplus residual; full construction in EC §EC.1.2), which holds if and only if  $F < F_{\text{crit}}$ . At  $F = F_{\text{crit}}$ , the surplus  $S^* = 0$  and the equilibrium ceases to exist. For  $F > F_{\text{crit}}$ ,  $G > 0$ , activating overload depletion of  $W$  through Eq. (6), which reduces  $h(W)$ , which further widens  $G$  through Loop R2. The reinforcing dynamics drive the system toward the low-trust attractor. Boundedness follows from Lemma 3 (EC §EC.1.2):  $\dot{K} \leq C(F) - \delta K$  by Gronwall's lemma, and  $W, T \in [0, 1]$  by construction. Full proof in EC §EC.1.3.  $\square$

At the calibrated baseline,  $F_{\text{crit}} \approx 1.11$ , corresponding to roughly a 40% increase in digital feature complexity above  $F_0 = 0.80$ . Figure 2 maps the regime structure: green cells converge to the surplus attractor (high CLV); red cells converge to the low-trust attractor ( $\text{CLV} \approx 400$ ). The cross-domain CLV sign reversal requires  $b > 1$  (super-linear demand) and  $p > 0$  (workforce-contingent threshold); at  $p = 0$  the threshold reduces to  $(u/\delta)^{1/b}$  and no longer responds to workforce investment.

<sup>1</sup> CLV is formally calibrated in §5; the theorems require only that CLV is increasing in  $T^*$ , which holds for any standard customer-lifetime-value functional.

**Figure 2** Basin-of-attraction map. Green cells converge to the surplus attractor (high CLV) for  $F < F_{\text{crit}} \approx 1.11$ ; red cells converge to the low-trust attractor for  $F \geq F_{\text{crit}}$ . Doubled-grid validation confirms boundary stability (EC §EC.4.2).



#### 4.3. Theorem 2: The Good Jobs Buffer

**THEOREM 2 (Good Jobs Buffer).** Suppose  $p > 0$  and the surplus equilibrium exists and is stable (Conditions C1–C4, i.e., the surplus equilibrium exists, is unique, and the low-trust attractor is present; Remark 1). Workforce investment  $f$  raises steady-state CLV through a multiplicative capacity-level channel:

$$\frac{\partial \text{CLV}^*}{\partial f} > 0 \quad \text{for all } f \geq 0. \quad (9)$$

The return decomposes into a direct capacity effect and an indirect trust effect:

$$\begin{aligned} \frac{\partial \text{CLV}^*}{\partial f} &= \underbrace{\frac{\partial \text{CLV}}{\partial T^*} \cdot \frac{\partial T^*}{\partial S^*} \cdot K^* \cdot p \cdot \frac{\partial W^*}{\partial f}}_{\text{direct: } h(W) \text{ raises } K_{\text{eff}}} \\ &+ \underbrace{\text{GE trust-demand feedback}}_{\text{indirect: } T^* \rightarrow D, \text{ dampened by } \approx 0.09; \text{ EC §EC.1.4}}. \end{aligned}$$

*Proof sketch.* Since  $W^* = (S^* + f)/(S^* + f + \mu)$  at the surplus equilibrium, differentiating gives  $\partial W^*/\partial f = \mu/(S^* + f + \mu)^2 > 0$ : workforce investment raises steady-state well-being. Higher  $W^*$  raises  $h(W^*) = (1 - p) + pW^*$ , which increases  $K_{\text{eff}} = K^* \cdot h(W^*)$ , which widens the surplus  $S^*$ , which raises  $T^* = \beta S^*/(\beta S^* + \lambda)$ , which raises CLV. Each link is strictly positive:  $\partial h/\partial W = p > 0$ ,  $\partial S/\partial h = K^* > 0$ ,  $\partial T^*/\partial S^* > 0$ ,  $\partial \text{CLV}/\partial T^* > 0$ . The chain-rule product is therefore positive, and diminishing returns follow from the concavity of  $W^*$  in  $f$ . Full proof in EC §EC.1.4.  $\square$

The  $p=0$  counterfactual in §5.1 confirms this: disabling the workforce multiplier leaves equilibrium CLV exactly invariant to workforce spending. The entire CLV return flows through service capacity, not the operational sensing channel: the indirect sensing channel contributes less than 1% of the total CLV uplift across all tested  $p$  values (EC Table EC.4). Standard dashboards, which record technology and workforce returns in separate ledgers, structurally underestimate the financial return to workforce investment.

#### 4.4. Theorem 3: The Burnout Cliff

**THEOREM 3 (Burnout Cliff).** *Suppose  $\sigma > 1$ ,  $p > 0$ , and  $F_0 < F_{\text{crit}}$  (a surplus equilibrium exists at  $F_0$ ). Under these conditions the system is bistable in a neighbourhood of  $F_{\text{crit}}$  (EC Remark 6). Consider a firm at the surplus equilibrium that implements a step increase in  $F$  from  $F_0$  to  $F_0 + \Delta F$ . Two thresholds govern the outcome:*

(i) *Permanent-collapse onset. For  $\Delta F \geq \Delta F_{\text{collapse}}$ , no surplus equilibrium exists at the post-upgrade level. The low-trust attractor is the only locally stable state (EC Remark 3); convergence to it is confirmed numerically across the studied initial conditions.*

$$\begin{aligned} \Delta F_{\text{collapse}} &= F_{\text{crit}} - F_0 \\ &= \left( \frac{u}{\delta} [(1-p) + p \frac{f}{f+\mu}] \right)^{1/b} - F_0. \end{aligned}$$

*For  $\Delta F < \Delta F_{\text{collapse}}$ , a new surplus equilibrium exists at the post-upgrade level. Whether a given trajectory reaches this equilibrium depends on its initial condition relative to the basin boundary (Figure 2); numerical verification confirms convergence from the pre-shock equilibrium initial condition across the parameter configurations studied (EC §EC.4.2).*

(ii) Instantaneous-overload diagnostic. A second, larger threshold marks the boundary above which the service gap turns positive at  $t = 0^+$ :

$$\Delta F_{\text{rapid}} = \frac{S_0}{b \cdot F_0^{b-1} \cdot (1 + aT^*)}, \quad (10)$$

where  $S_0 = K^*h(W^*) - D(T^*, F_0)$  is the pre-upgrade surplus. Because  $\Delta F_{\text{rapid}} > \Delta F_{\text{collapse}}$  (at baseline: 0.40 vs. 0.31), a firm in the zone  $\Delta F \in [\Delta F_{\text{collapse}}, \Delta F_{\text{rapid}})$  collapses permanently but through a slow trajectory ( $G$  negative at  $t = 0^+$ , surplus eroded gradually). For  $\Delta F \geq \Delta F_{\text{rapid}}$ ,  $G$  turns positive immediately and the system enters a rapid self-reinforcing collapse to the low-trust attractor.

The safe zone is  $\Delta F < \Delta F_{\text{collapse}}$ ;  $\Delta F_{\text{rapid}}$  is a deployment-speed diagnostic, not a safety threshold.

*Proof sketch.* Part (i) follows directly from Theorem 1: for  $F_0 + \Delta F \geq F_{\text{crit}}$ , no surplus equilibrium exists; for  $\Delta F < \Delta F_{\text{collapse}}$ , a new surplus equilibrium exists at  $F_0 + \Delta F$ , and numerical verification confirms that trajectories originating at the pre-shock surplus equilibrium converge to this new surplus equilibrium across the parameter configurations studied (EC §EC.4.2). For Part (ii), setting the first-order demand increment  $\Delta D \approx (1 + aT^*) b F_0^{b-1} \Delta F$  equal to the pre-upgrade surplus  $S_0$  gives  $\Delta F_{\text{rapid}}$ . When  $\Delta F < \Delta F_{\text{rapid}}$ ,  $G_0^+ = 0$  at  $t_0^+$  and the system crosses the workforce tipping threshold through a slow transient, consistent with the slow-collapse zone. Full proof in EC §EC.1.5.  $\square$

Workforce investment  $f$  raises  $F_{\text{crit}}$  and enlarges the safe upgrade zone. Before each deployment phase, the firm should compute  $\Delta F_{\text{collapse}} = F_{\text{crit}}(f) - F_0$  as a hard ceiling on planned upgrade size; this ceiling is set entirely by  $f$ .

#### 4.5. Proposition 4: Digitalisation-Workforce Complementarity

**PROPOSITION 4 (Digitalisation-Workforce Complementarity).** *The marginal value of workforce improvement is strictly increasing in the technology base:*

$$\begin{aligned} \frac{\partial K_{\text{eff}}}{\partial W} &= p \cdot K, \\ \frac{\partial^2 K_{\text{eff}}}{\partial W \partial K} &= p > 0. \end{aligned} \quad (11)$$

Since maintaining surplus at higher  $F$  requires higher effective capability  $K_{\text{eff}}$  (that is, either higher  $K^*$  through increased  $u$ , or higher  $h(W^*)$  through increased  $f$ ), the equilibrium return to workforce investment increases with digitalisation level.



A practical corollary follows: workforce well-being sets a ceiling on how far digitalisation can safely go. At the regime boundary ( $S^* \rightarrow 0$ ), the equilibrium trust approaches zero ( $T^* \rightarrow 0$ ; see Lemma 2, EC §EC.1.2), so:

$$F_{\max}(W) = (K^* \cdot h(W))^{1/b}, \quad (12)$$

which is strictly increasing in  $W$  for  $p > 0$ . The full implicit differentiation accounting for  $T^*(W, F)$  dependence at interior equilibria (where  $S^* > 0$ ) yields a correction term of order  $O(aS^*)$  that vanishes as  $S^* \rightarrow 0$ ; the formula (12) is therefore exact at the boundary and an upper bound on  $F_{\max}$  in the interior: trust-elevated demand means true safe capacity is approximately 8% lower than this formula implies at baseline trust levels (EC §EC.1.6).

*Proof.* The cross-partial follows directly from  $K_{\text{eff}} = K[(1-p) + pW]$ . The  $F_{\max}$  formula identifies  $F_{\text{crit}}$  at  $S^* \rightarrow 0$ , where  $T^* \rightarrow 0$ , giving  $(K^* h(W))^{1/b}$ , strictly increasing in  $W$  since  $h'(W) = p > 0$ ; with the full  $T^*(W, F)$  correction this is exact at the boundary and an upper bound in the interior (EC §EC.1.6).  $\square$


*Remark.* The constant cross-partial  $\partial^2 K_{\text{eff}} / \partial W \partial K = p$  is an immediate consequence of the linear specification  $h(W) = (1-p) + pW$ . Under a concave alternative  $h(W) = (1-p) + pW^\gamma$  ( $\gamma < 1$ ), the cross-partial becomes  $p\gamma W^{\gamma-1}$ , which is state-dependent: complementarity is stronger when the workforce is depleted (low  $W$ ), precisely when it matters most operationally. The linear form therefore understates the complementarity effect under overload, making our result a conservative bound. Section 7.2 discusses additional implications of concave  $h$ .

Figure 3 illustrates: the gap between a stressed firm ( $W = 0.30$ ,  $F_{\max} \approx 1.063$ ) and a healthy one ( $W = 0.70$ ,  $F_{\max} \approx 1.307$ ) is 23%, arising entirely from workforce investment.

**COROLLARY 5 (Budget-Constrained Trap).** *Under a fixed budget constraint  $c_F F + c_f f \leq B$  with  $B \geq B_{\min} - \varepsilon$  (where  $B_{\min} = c_F \cdot F_{\text{crit}}$ ), the CLV-maximising allocation includes positive workforce investment  $f^* > 0$ : the all- $F$  corner is strictly dominated analytically whenever the budget reaches or exceeds the Trap threshold, and computationally for  $B$  throughout  $[B_{\min}/2, B_{\min})$  across cost ratios  $c_F/c_f \in [0.5, 5]$ .*

*Proof.* See EC §EC.1.7.  $\square$

**Figure 3** Digitalisation-workforce complementarity. The healthy firm ( $W = 0.70$ ) sustains 23% higher  $F_{\max}$  than the stressed firm ( $W = 0.30$ ), arising entirely from workforce investment.




## 5. Computational Study

Numbers sharpen the intuition. We calibrate the model to baseline parameters (Table 2) and simulate three deployment strategies that span the critical threshold. CLV follows Gupta et al. (2004):  $CLV = w(T)/(1 - \rho(T))$ , where  $\rho(T)$  is an annual retention probability and  $w(T)$  an annualised per-customer margin, with  $\rho(T) = 0.80 + 0.18T$  and  $w(T) = 80(1 + 0.25 \max(T - 0.15, 0))$ . The endpoints  $\rho(0) = 0.80$  and  $\rho(1) = 0.98$  fall within the 70–90% annual range in Gupta et al. (2004); \$80 is an illustrative annual margin for US omnichannel grocery. All analytical results hold for any CLV increasing in trust (calibration in EC §EC.4.3). All three qualitative phenomena are robust across  $b \in [1.1, 2.0]$ ,  $\sigma \in (1.0, 3.0]$ , and  $p \in [0.25, 0.75]$  (EC Table EC.3; at  $\sigma = 1$  the Burnout Cliff disappears).

### 5.1. The Hidden Cost of App Innovation (Figure 4)

Three deployment strategies span the critical threshold  $F_{\text{crit}} \approx 1.11$ . All begin from  $(K_0, W_0, T_0) = (1.60, 0.55, 0.50)$  with  $F_0 = 0.80$ . Scenario A ramps aggressively to  $F_{\text{end}} = 1.15$  with no workforce investment; Scenario B ramps to  $F_{\text{end}} = 1.05$  with concurrent support ( $f = 0.06$ ); Scenario C pre-invests in workforce ( $f = 0.10$ , months 1–12, after which  $f$  returns to 0.0) before ramping to  $F_{\text{end}} = 1.15$ .

**Figure 4** The hidden cost of app innovation. Panel (a): the demand-volume index  $F \cdot D(T(t), F(t))$  rises for all three scenarios, masking divergent back-stage outcomes. Panel (b): CLV diverges sharply by month 24. Panels (c,d): workforce collapse precedes trust collapse in the no-investment scenario.




All three appear profitable on front-stage metrics (panel a), yet CLV diverges sharply (panel b): by month 24, Scenario C sustains CLV approximately 40% above Scenario A. Scenario C's persistent advantage—despite  $f$  returning to zero after month 12—confirms that well-being is a durable stock. Workforce collapse precedes trust collapse (panels c, d), confirming the invisible cross-domain cost.

## 5.2. Basin-of-Attraction Structure

Figure 2 maps the basin boundary across the joint space of  $W_0$  and  $F$ . The boundary shifts rightward as  $W_0$  increases: a firm near the basin boundary can increase its safety margin by raising workforce investment  $f$ , which shifts  $F_{\text{crit}}$  rightward without touching the technology stock.

**Figure 5 The Burnout Cliff. Normalised terminal CLV versus upgrade size  $\Delta F$  for three workforce-investment levels. Dotted verticals:  $\Delta F_{\text{collapse}}$  (0.18, 0.31, 0.49). The neglected firm ( $f = 0.01$ ) exhausts its safe zone at  $\Delta F = 0.18$ ; the invested firm ( $f = 0.10$ ) retains headroom to  $\Delta F = 0.49$ .**



### 5.3. The Burnout Cliff (Figure 5)

We sweep  $\Delta F \in [0, 0.60]$  for three workforce-investment scenarios: neglected ( $f = 0.01$ ), baseline ( $f = 0.03$ ), and invested ( $f = 0.10$ ), each starting from its surplus equilibrium at  $F_0 = 0.80$ . Varying  $f$  shifts  $F_{\text{crit}}(f)$  from 0.98 to 1.29, widening  $\Delta F_{\text{collapse}}$  from 0.18 to 0.49.

Figure 5 shows the three normalised CLV curves. For the neglected firm ( $f = 0.01$ ), the safe zone collapses to  $\Delta F < 0.18$ ; for the invested firm ( $f = 0.10$ ), it extends to  $\Delta F = 0.49$ . The slow-collapse zone between  $\Delta F_{\text{collapse}}$  and  $\Delta F_{\text{rapid}}$  is the diagnostic blind spot: dashboards report continued growth while trust has already turned negative. For the well-investing firm ( $f = 0.10$ ), the ordering reverses— $\Delta F_{\text{collapse}} = 0.49 > \Delta F_{\text{rapid}} \approx 0.40$ —so the slow-collapse zone disappears entirely (EC §EC.1.5): a well-investing firm faces either a fully safe zone or immediately visible overload. Workforce investment thus purchases perfect dashboard visibility of deployment risk, making invisible gradual erosion the exclusive hazard of underinvesting firms.

**Table 3** Model-to-data mapping.

| Model variable             | Empirical proxy                   | Justification                      |
|----------------------------|-----------------------------------|------------------------------------|
| $F$ (digital depth)        | E-commerce share (FRED ECOMPCTSA) | National digital penetration index |
| $W$ (workforce well-being) | Glassdoor overall rating (1–5)    | Composite employee engagement      |
| $T$ (customer trust)       | ACSI score (0–100)                | Aggregate customer confidence      |
| $K$ (technology stock)     | CapEx / Revenue                   | Back-stage capital intensity       |

## 6. Empirical Evidence

The model makes testable predictions at two levels of aggregation. At the firm level, we test the Front-Stage Trap using an inverted-U specification ( $N = 216$ , 17 firms) and an event study around major omnichannel launches ( $N = 108$ , 16 firms<sup>2</sup>). A cross-sectional complementarity check (Glassdoor  $\times$  technology interaction) provides a directional sign check but lacks statistical power ( $N_e = 13$ ). At the county level, we test a labour-market analogue of the Good Jobs Buffer: whether counties with higher initial retail wages show more resilient employment as e-commerce grows ( $N = 30,951$  county–year observations, 3,115 US counties). An independent European sample ( $N_e = 11$ ,  $N = 20$ ) is reported in EC §EC.5 as an exploratory check.

### 6.1. Data and Samples

Table 3 maps each model variable to its empirical proxy.

*Firm-level sample.* We assemble an annual panel of 17 publicly listed US retailers from 2008 to 2024 using SEC EDGAR 10-K filings. Customer satisfaction ( $T$  proxy) is measured by the American Customer Satisfaction Index (ACSI; Fornell et al. 1996;  $N = 216$ , 17 firms). Workforce well-being ( $W$  proxy) uses Glassdoor overall ratings (1–5 scale; restricted to firm–years with at least five reviews), a publicly available aggregate of employee reviews that has been validated as a measure of employee satisfaction and used as a proxy in employee-wellbeing research (Edmans 2011). Technology investment ( $K$  proxy) is capital expenditure as a share of revenue. E-commerce share is from the Federal Reserve (FRED series ECOMPCTSA). The inverted-U specification requires only ACSI and e-commerce ( $N = 216$ , 17 firms); the complementarity specification additionally requires Glassdoor and CapEx, yielding 131 complete firm–year observations across 13 firms after listwise deletion.

<sup>2</sup> Amazon is excluded from the event study because its omnichannel transition was gradual and continuous rather than marked by a discrete launch event, making event-time identification infeasible.

**Table 4** Descriptive statistics.

| Variable                                                            | N      | Mean                       | SD     | Range      |
|---------------------------------------------------------------------|--------|----------------------------|--------|------------|
| <i>Panel A: Firm-level sample (13 firms, 2010–2024<sup>‡</sup>)</i> |        |                            |        |            |
| ACSI score                                                          | 131    | 76.5                       | 3.4    | 68–88      |
| Glassdoor rating (1–5)                                              | 131    | 3.23                       | 0.30   | 2.7–3.9    |
| CapEx / Revenue (%)                                                 | 131    | 2.9 [2.0–4.5] <sup>†</sup> |        |            |
| <i>Panel B: County-level sample (3,115 counties, 2015–2024)</i>     |        |                            |        |            |
| Annual avg. retail employment                                       | 30,951 | 5,035                      | 15,458 | 18–423,900 |
| Annual avg. weekly wage (\$)                                        | 30,951 | 558                        | 151    | 203–2,861  |
| Baseline wage, 2015 (\$)                                            | 3,115  | 476                        | 110    | 203–1,697  |
| Baseline employment, 2015                                           | 3,115  | 5,053                      | 15,712 | 51–416,188 |
| Log employment growth (vs. 2015)                                    | 30,951 | 0.00                       | 0.12   | –1.5–2.9   |
| National e-commerce share (%)                                       | 10     | 11.8                       | 3.0    | 7.4–16.1   |

<sup>†</sup>CapEx/Revenue reports the median [IQR] rather than the mean and SD due to extreme outliers in the raw ratio (e.g., BBY 2012). Panel A describes the 13-firm subsample with joint ACSI, Glassdoor, and CapEx coverage (used for the complementarity test). <sup>‡</sup>Panel A starts in 2010 because systematic Glassdoor coverage begins circa 2010; earlier observations are excluded by listwise deletion. The inverted-U and event-study specifications use a broader 17-firm and 16-firm ACSI-only sample ( $N = 216$  and  $N = 108$ , respectively); descriptive statistics for these samples are available upon request.

*County-level sample.* The BLS Quarterly Census of Employment and Wages provides county-level annual average retail employment, weekly wages, and establishment counts for the retail trade sector (NAICS 44–45), 2015–2024. We retain counties with baseline retail employment exceeding 50 and weekly wage exceeding \$200, yielding 3,115 counties and 30,951 county–year observations (199 of the theoretical 31,150 are dropped due to missing QCEW values).

Table 4 reports descriptive statistics for both samples.

## 6.2. Firm-Level Results

*Cross-sectional complementarity (Proposition 4).* The between-effects (OLS) estimate of the standardised interaction  $K^z \times W^z$  is positive (+1.95,  $p = 0.41$ ;  $N = 131$ , 13 firms). The sign is consistent with the prediction that firms with jointly higher technology and workforce investment enjoy higher satisfaction, and aligns directionally with Aral et al. (2012). However, the coefficient is not statistically significant, reflecting the extreme power limitation of 13 firm clusters (Cameron et al. 2008). We treat this as a sign check, not a test. The within-firm TWFE estimate of the interaction is negative and marginally significant ( $-0.56^*$ ,  $p < 0.10$ ). This does not reflect the structural relationship: within-firm Glassdoor SD is only 30% of total

SD, so in years when a firm simultaneously raises CapEx and satisfaction, the interaction captures short-run managerial alignment (common-cause co-movement) rather than structural cross-firm complementarity. The negative main effect of  $K^z$  in the TWFE ( $-0.33^{**}$ ) is consistent with the Front-Stage Trap: within firm, years of higher CapEx intensity are associated with lower satisfaction, reflecting operational disruption preceding capability gains. The between-effects estimand, exploiting persistent cross-firm differences, is the appropriate estimand for the structural question (Milgrom and Roberts 1990).

*Inverted-U (Theorem 1).* We estimate the quadratic specification

$$\text{ACSI}_{it} = \alpha_i + \beta_1 \text{Ecom}_t + \beta_2 \text{Ecom}_t^2 + \mathbf{X}_{it}\theta + \varepsilon_{it}, \quad (13)$$

where  $\alpha_i$  are firm fixed effects,  $\text{Ecom}_t$  is national e-commerce share, and  $\mathbf{X}_{it}$  includes workforce controls where indicated. Year fixed effects are excluded because  $\text{Ecom}_t$  is a national aggregate identical within year; including them would absorb it entirely. Identification exploits the time-series curvature of the national digitalisation trend, shared across all 17 firms. Firm-specific deployment timing would sharpen identification; the event study provides this. As a supplementary check, replacing  $\text{Ecom}_t$  with a firm-specific digitalization index (cumulative technology adoption percentage;  $N = 84$ , 6 firms with joint coverage) yields a significant inverted-U ( $\hat{\beta}_2 = -0.00061$ ,  $p = 0.039$ , firm FE; turning point at 51% adoption), confirming the non-monotonicity with firm-level variation.

The specification ( $N = 216$ , 17 firms) yields  $\hat{\beta}_2 = -0.028$  ( $p = 0.028$  under  $t(16)$ ), consistent with the predicted inverted-U. The estimated turning point (7.9% e-commerce share) lies in the mid-2010s, qualitatively consistent with a digitalisation ceiling, though the correspondence is directional, not quantitative.

The coefficient attenuates when workforce controls are included ( $N = 190$ ;  $p = 0.66$ ). A Baron-Kenny mediation analysis confirms the mechanism: Glassdoor mediates approximately 81% of the total e-commerce–ACSI relationship, and the direct effect becomes insignificant ( $p = 0.70$ ) once workforce wellbeing is controlled. This is consistent with the model's core prediction that digitalisation damages satisfaction *through* workforce depletion, not through a direct channel.

*Event study (Theorem 1).* We exploit the staggered timing of major omnichannel launches (e.g., Walmart curbside 2017, Target Drive Up 2018, Kroger ClickList 2016) as firm-level technology shocks. Using a

$[-3, +3]$  year event window around each firm's launch date  $t_i^*$  ( $N = 108$ , 16 firms), the event-study specification is

$$\begin{aligned} \text{ACSI}_{it} = & \alpha_i + \gamma_t + \sum_{\tau=-3}^{+3} \delta_\tau \cdot \mathbf{1}[t - t_i^* = \tau] \\ & + \mathbf{X}_{it}\theta + \varepsilon_{it}, \end{aligned} \quad (14)$$

where  $\alpha_i$  are firm fixed effects,  $\gamma_t$  are year fixed effects,  $\tau = t - t_i^*$  is event time,  $\delta_{-1} = 0$  is the omitted reference period, and  $\mathbf{X}_{it}$  includes pre-shock capital expenditure as a control. The averaged post-shock coefficient ( $\bar{\delta}_{+1,+3}$ ) yields an ACSI decline of  $-0.71$  ( $p = 0.037$  under  $t(15)$ );  $t = 0$  is excluded because it captures transition noise (partial rollout, announcement effects); the  $t = 0$  coefficient is itself significant ( $-0.37$ ,  $p = 0.043$ ; EC Table EC.6). The result is consistent with the Front-Stage Trap. Event-time coefficients for the level specification show near-zero and insignificant pre-trend coefficients at  $t = -3$  and  $t = -2$  (relative to  $t = -1$ ), supporting the parallel-trends assumption for the basic difference-in-differences (EC Table EC.6). Pre-trend coefficients for the workforce-moderation interaction are similarly insignificant (EC Table EC.6). The workforce moderation ( $\text{post} \times \text{pre-Glassdoor}$ ) is not detectable ( $p = 0.79$ ), consistent with the low within-firm Glassdoor variation rather than absence of the mechanism.

A caveat on staggered treatment timing: recent work (Callaway and Sant'Anna 2021, Sun and Abraham 2021) shows that TWFE estimators with heterogeneous staggered treatment can be biased. Two features limit this concern here: the sample is small ( $N_c = 16$  firms) with launches concentrated in 2016–2019, limiting cohort heterogeneity; and the target estimand is the average post-shock effect, which TWFE recovers consistently under modest heterogeneity. With only 16 treated firms, cohort-specific ATTs from Callaway-Sant'Anna or Sun-Abraham would be highly imprecise; we rely on standard TWFE under our maintained assumption of limited treatment-effect heterogeneity across 2016–2019 launch cohorts, supported by the pre-trend evidence.

### 6.3. County-Level Good Jobs Buffer

The firm-level data cannot deliver a well-powered test of the capacity-multiplier channel. We therefore test an *analogue* of the Good Jobs Buffer at a different margin (employment resilience rather than customer satisfaction) using BLS county-level data, where the dramatically larger sample ( $N = 30,951$ ) provides the statistical power that the firm-level analysis lacks.



*Design.* The model predicts that workforce investment  $f$  buffers the negative effects of digitalisation on service capacity (Theorem 2). The county-level test maps this prediction to an observable analogue: do counties with higher initial retail wages (a proxy for local workforce-investment conditions) show more resilient retail employment as national e-commerce penetration grows? Three mapping choices warrant acknowledgment: the outcome is employment resilience (not CLV); county wages are a noisy proxy for individual-firm  $f$  (reflecting local labour-market conditions including minimum wages and cost of living); and  $Ecom_t$  is a national macro trend whereas  $F$  is a firm-level choice. Despite these distinctions, the qualitative prediction—higher workforce investment buffers against digitalisation-induced capacity loss—is preserved. Initial retail wage is measured in 2015 (before the e-commerce acceleration), standardised cross-sectionally. The interaction  $W_{c,2015}^z \times Ecom_t$  captures whether high-wage counties differentially maintain employment under digital pressure:

$$\begin{aligned} \ln E_{ct} - \ln E_{c,2015} &= \alpha_c + \gamma_t \\ &+ \beta_1 (W_{c,2015}^z \times Ecom_t) + \varepsilon_{ct}. \end{aligned} \quad (15)$$

Because  $Ecom_t$  is a national-level variable (identical across counties within each year), its main effect is perfectly collinear with the year fixed effects  $\gamma_t$  and is therefore absorbed by them. Likewise,  $W_{c,2015}^z$  is time-invariant and absorbed by the county fixed effects  $\alpha_c$ . The coefficient  $\beta_1$  thus identifies the *differential* effect of e-commerce growth on counties with higher initial retail wages, which is the quantity of interest for the Good Jobs Buffer.

*Results.* The interaction is positive and significant ( $\hat{\beta}_1 = 0.0026$ ,  $t = 2.24$ ,  $p = 0.025$ ; county and year FE, clustered at county). A county one SD above the mean in initial retail wage ( $\approx \$110/\text{week}$ ) retains 2.6 percentage points more retail employment per 10-point increase in e-commerce share, consistent with the Good Jobs Buffer operating through employment capacity. The long-difference specification (2015 vs. 2023) confirms: high-wage counties gained 2.1 pp more retail employment over the period ( $t = 2.57$ ,  $p = 0.010$ ).

*Robustness.* The result strengthens when COVID years are excluded ( $\hat{\beta}_1 = 0.0031$ ,  $p = 0.014$ ), holds in urban counties ( $p = 0.006$ ), and is insignificant in rural counties ( $p = 0.58$ ). A pre-period placebo (2015–2017) shows no effect ( $p = 0.40$ ), confirming the buffer activates only under substantial digital pressure. The

**Table 5 Firm-level regression results.**

|                       | <i>Complementarity</i> |                   | <i>Inverted-U</i>   |                   |
|-----------------------|------------------------|-------------------|---------------------|-------------------|
|                       | (1) OLS                | (2) TWFE          | (3) Firm FE         | (4) + $W$         |
| $K^z$                 | -0.92<br>(0.73)        | -0.33**<br>(0.13) |                     |                   |
| $W^z$                 | 2.84***<br>(0.51)      | 2.00**<br>(0.68)  |                     |                   |
| $K^z \times W^z$      | 1.95<br>(2.36)         | -0.56*<br>(0.31)  |                     |                   |
| Ecom                  |                        |                   | 0.44*<br>(0.23)     | 0.03<br>(0.27)    |
| Ecom <sup>2</sup>     |                        |                   | -0.028**<br>(0.012) | -0.005<br>(0.012) |
| Glassdoor             |                        |                   |                     | 9.87***<br>(1.80) |
| Firm FE               | No                     | Yes               | Yes                 | Yes               |
| Year FE               | No                     | Yes               | No                  | No                |
| $N$                   | 131                    | 131               | 216                 | 190               |
| Firms                 | 13                     | 13                | 17                  | 17                |
| $R^2_{\text{within}}$ | —                      | 0.44              | 0.18                | 0.51              |

Dependent variable: ACSI score. Cluster-robust standard errors (at firm level) in parentheses;  $p$ -values use  $t(G-1)$  degrees of freedom. \*\*\* $p < 0.01$ ; \*\* $p < 0.05$ ; \* $p < 0.10$ .  $K^z$  and  $W^z$  are standardised CapEx/Revenue and Glassdoor overall, respectively. Ecom is national e-commerce share (%). Columns (1)–(2): complementarity specification ( $N=131$ , 13 firms). Columns (3)–(4): inverted-U specification ( $N=216/190$ , 17 firms). Turning point for column (3): 7.9%. Attenuation of Ecom<sup>2</sup> in column (4) is consistent with the model's prediction that e-commerce damages satisfaction through workforce depletion. The negative within-firm TWFE estimate in column (2) reflects low temporal variation in Glassdoor ratings (within-firm SD is 30% of total SD); see text. Column (1) is OLS without fixed effects;  $R^2_{\text{within}}$  is not applicable (—); overall  $R^2$  for column (1): 0.55.

effect operates through employment quantity, not wages (wage-growth outcome:  $p = 0.98$ ), consistent with the capacity-multiplier channel.

A COVID caveat: the e-commerce share jump from 10.7% (2019) to 14.0% (2020) represents forced adoption rather than strategic choice. The result's robustness when COVID years are excluded ( $p = 0.014$ ) confirms the effect is not driven solely by pandemic-era variation.

Tables 5 and 6 present the full regression results.

#### 6.4. Assessment

Three empirical strategies support the model's qualitative predictions. The inverted-U ( $p = 0.028$ ) and the post-shock ACSI decline ( $p = 0.037$ ) both support the Front-Stage Trap (Theorem 1). The county-level wage–e-commerce interaction ( $p = 0.025$ ,  $N = 30,951$ ) provides a well-powered test of the Good Jobs Buffer (Theorem 2); the low within- $R^2$  (0.008) is expected for a single interaction net of 3,115 county and

**Table 6 County-level regression results: Good Jobs Buffer.**

|                                     | (1) Main             | (2) Excl. COVID      | (3) Urban             | (4) Rural          | (5) Placebo        |
|-------------------------------------|----------------------|----------------------|-----------------------|--------------------|--------------------|
| $W_{c,2015}^z \times \text{Ecom}_t$ | 0.0026**<br>(0.0011) | 0.0031**<br>(0.0013) | 0.0050***<br>(0.0018) | 0.0008<br>(0.0014) | 0.0023<br>(0.0027) |
| County FE                           | Yes                  | Yes                  | Yes                   | Yes                | Yes                |
| Year FE                             | Yes                  | Yes                  | Yes                   | Yes                | Yes                |
| $N$                                 | 30,951               | 24,768               | 15,532                | 15,419             | 9,321              |
| Counties                            | 3,115                | 3,115                | 1,557                 | 1,558              | 3,115              |
| $R^2_{\text{within}}$               | 0.008                | 0.013                | 0.042                 | −0.000             | 0.001              |

Dependent variable:  $\ln E_{ct} - \ln E_{c,2015}$  (log employment growth relative to 2015 baseline). Standard errors clustered at county level in parentheses. \*\*\* $p < 0.01$ ; \*\* $p < 0.05$ .  $\text{Ecom}_t$  main effect absorbed by year FE (national variable);  $W_{c,2015}^z$  absorbed by county FE (time-invariant). Column (2): excludes 2020–2021. Columns (3)–(4): split at median baseline employment. Column (5): 2015–2017 only (pre-period placebo; 24 county-year observations dropped due to missing BLS data in 2015–2017). Negative within- $R^2$  in column (4) indicates the interaction term does not improve fit beyond fixed effects in the rural subsample, consistent with the buffer mechanism operating primarily in urban retail markets.

10 year fixed effects, comparable to Chetty et al. (2009). The complementarity sign (+1.95,  $p = 0.41$ ) is directionally consistent with Proposition 4 but underpowered. The contribution is directionally consistent evidence across different identification strategies, levels, and outcomes, subject to §7.2.

## 7. Discussion

### 7.1. Contributions and Managerial Implications

The three theorems together provide the formal micro-foundation that connects Ton’s (2014) four operational choices to each other through the dynamics of  $K_{\text{eff}}$ . Each mapping is analogical—Ton’s prescriptions address firm-level operational choices whereas the model’s parameters index structural mechanisms—but the correspondences are informative. “Operate with slack” corresponds to positive  $f$ ; Theorem 2 quantifies the CLV return. “Standardise and empower” corresponds to lower  $\sigma$ ; well-designed processes enlarge the safe zone of Theorem 3. “Offer less” maps onto keeping  $F$  below  $F_{\text{crit}}$ . “Cross-train” maps onto higher  $p$ : cross-trained workers raise  $h(W)$  at any engagement level (Proposition 4). The model extends Oliva–Stermann with customer trust, a front-/back-stage domain split, and digital scaling: three extensions absent from the quality-erosion framework, now formalised in a single model that, to the best of our knowledge, provides the first dynamic formalisation of the service profit chain (Heskett et al. 1994, 1997) and Good Jobs Strategy (Ton 2014) with computable boundary conditions.

**Managerial implications.** Four implications follow. First, the maximum safe rollout speed is a workforce constraint (Theorem 3): the firm should verify that  $\Delta F_{\text{collapse}} = F_{\text{crit}}(f) - F_0$  exceeds the planned upgrade

size. Second, workforce well-being is a leading indicator:  $W$  reaches crisis levels 6–21 months before trust does (EC Table EC.4). Third, the more digital the firm becomes, the more it should invest in its people (Proposition 4): conventional ROI calculations underestimate returns because the dominant capacity-multiplier channel is invisible in standard dashboards (Theorem 2). Fourth, as a direct consequence of Theorems 1 and 2, the all-technology allocation is dominated at and near the Trap threshold (analytically) and throughout the tested budget range (computationally): reallocating the marginal dollar to workforce spending yields a strict CLV gain (Corollary 5).

## 7.2. Limitations and Directions for Future Research

*Model limitations.*  $F$  is exogenous; the HJB formulation for optimal deployment is non-trivial because the switching manifold  $\{G=0\}$  is endogenous (Filippov 1988). Demand grows without bound in  $F$  (no market saturation); predictions are most reliable near  $F_{\text{crit}}$ . Omitting a direct  $W \rightarrow D$  channel and using linear  $h(W)$  both make results conservative; a concave  $h(W) = (1-p) + pW^\gamma$  ( $\gamma < 1$ ), motivated by prospect theory's diminishing sensitivity, would strengthen them.

*Empirical limitations.* With  $N_c = 13$  firm clusters, firm-level power is limited (Cameron et al. 2008); county-level power ( $N_c = 3,115$ ) does not share this limitation. Glassdoor's low within-firm variation (SD 30% of total) limits the interaction's within-firm identification. The county interaction cannot fully separate the Good Jobs Buffer from non-tradeable-retail embedding; store-level data would resolve both.

*Extensions.* Store-level identification of  $p$ ; competitive multi-firm extension with workforce poaching; stochastic demand forcing (probabilistic Burnout Cliff).

## 8. Conclusion

Standard operational dashboards measure the front-stage domain that invests in technology, not the back-stage workforce that fulfils the resulting demand and absorbs the consequences. This disconnect is not merely a measurement problem; it is a structural feature of omnichannel operations that makes cross-domain value destruction structurally predictable from two observable managerial choices: the workforce investment rate  $f$  and the digitalisation level  $F$ . This paper provides the formal instrument for both.

Together, the three theorems establish that the relationship between digitalisation and customer value is non-monotonic (the Front-Stage Trap imposes a computable ceiling), workforce-contingent (the Good

Jobs Buffer channels the entire return through service capacity), and bistable (the Burnout Cliff creates a deployment threshold that workforce investment widens by a factor of 2.7 at baseline). These three properties are invisible to dashboards that partition technology costs and workforce costs onto separate ledgers, which is precisely why the structural gap persists.

These results give operations teams two computable numbers standard dashboards do not report: a digitalisation ceiling ( $F_{\text{crit}}$ ) and a maximum safe rollout size ( $\Delta F_{\text{collapse}}$ ). Three empirical strategies provide directionally consistent evidence across different levels of analysis and identification strategies: an inverted-U between e-commerce and ACSI, a post-shock satisfaction decline following omnichannel launches, and county-level evidence that higher retail wages buffer employment against e-commerce growth ( $N = 30,951$  county-year observations, 3,115 counties,  $p = 0.025$ ). Store-level identification of  $p$  using exogenous technology rollout events remains the most important empirical follow-up. The broader implication extends beyond omnichannel retail: wherever automated front-stage processes generate demand on human back-stage operations, workforce investment is an operational prerequisite for digital strategy, not an optional cost line. The capacity-multiplier structure is plausibly applicable wherever technology adoption generates hidden back-stage load—healthcare, financial services, logistics—though whether the quantitative thresholds apply in these domains requires domain-specific calibration and is a direction for future work.

## References

- Acimovic J, Graves SC (2015) Making better fulfillment decisions on the fly in an online retail environment. *Manufacturing & Service Operations Management* 17(1):34–51.
- Aral S, Brynjolfsson E, Wu L (2012) Three-way complementarities: Performance pay, human resource analytics, and information technology. *Management Science* 58(5):913–931.
- Argote L, Epple D (1990) Learning curves in manufacturing. *Science* 247(4945):920–924.
- Bainbridge L (1983) Ironies of automation. *Automatica* 19(6):775–779.
- Bakker AB, Demerouti E (2007) The Job Demands-Resources model: State of the art. *Journal of Managerial Psychology* 22(3):309–328.
- Bell DR, Gallino S, Moreno A (2018) Offline showrooms in omnichannel retail: Demand and operational benefits. *Management Science* 64(4):1629–1651.

- Bijmolt THA, Broekhuis M, de Leeuw S, Hirche C, Rooderkerk RP, Sousa R, Zhu SX (2021) Challenges at the marketing-operations interface in omni-channel retail environments. *Journal of Business Research* 122:864–874.
- Callaway B, Sant'Anna PHC (2021) Difference-in-differences with multiple time periods. *Journal of Econometrics* 225(2):200–230.
- Cameron AC, Gelbach JB, Miller DL (2008) Bootstrap-based improvements for inference with clustered errors. *Review of Economics and Statistics* 90(3):414–427.
- Chetty R, Looney A, Kroft K (2009) Salience and taxation: Theory and evidence. *American Economic Review* 99(4):1145–1177.
- Corbett CJ (2024) OM Forum: The operations of well-being: An operational take on happiness, equity, and sustainability. *Manufacturing & Service Operations Management* 26(2):409–430.
- Demerouti E, Bakker AB, Nachreiner F, Schaufeli WB (2001) The Job Demands-Resources model of burnout. *Journal of Applied Psychology* 86(3):499–512.
- Dubé A, Reich A, Shierholz H, West M (2020) Minimum wage shocks, employment flows and labor market frictions. *Journal of Labor Economics* 38(4):1009–1052.
- Edmans A (2011) Does the stock market fully value intangibles? Employee satisfaction and equity prices. *Journal of Financial Economics* 101(3):621–640.
- Filippov AF (1988) *Differential Equations with Discontinuous Righthand Sides*. Dordrecht: Kluwer Academic Publishers.
- Fisher ML, Gallino S, Netessine S (2021) Setting retail staffing levels: A methodology validated with implementation. *Manufacturing & Service Operations Management* 23(6):1562–1579.
- Fornell C, Johnson MD, Anderson EW, Cha J, Bryant BE (1996) The American Customer Satisfaction Index: Nature, purpose, and findings. *Journal of Marketing* 60(4):7–18.
- Gans JS, Goldfarb A (2026) O-ring automation. NBER Working Paper No. 34639. National Bureau of Economic Research, Cambridge, MA. Accessed March 2026.
- Gupta S, Lehmann DR, Stuart JA (2004) Valuing customers. *Journal of Marketing Research* 41(1):7–18.

- Heskett JL, Jones TO, Loveman GW, Sasser Jr WE, Schlesinger LA (1994) Putting the service-profit chain to work. *Harvard Business Review* 72(2):164–174.
- Heskett JL, Sasser Jr WE, Schlesinger LA (1997) *The Service Profit Chain: How Leading Companies Link Profit and Growth to Loyalty, Satisfaction, and Value*. New York: Free Press.
- Kahneman D, Tversky A (1979) Prospect theory: An analysis of decision under risk. *Econometrica* 47(2):263–291.
- Hübner A, Amorim P, Fransoo JC, Honhon D, Kuhn H, Martinez de Albeniz V, Robb D (2021) Digitalization and omnichannel retailing: Innovative OR approaches for retail operations. *European Journal of Operational Research* 294(3):817–819.
- KC D, Terwiesch C (2009) Impact of workload on service time and patient safety: An econometric analysis of hospital operations. *Management Science* 55(9):1486–1498.
- Kesavan S, Lambert SJ, Williams JC, Pendem PK (2022) Doing well by doing good: Improving retail store performance with responsible scheduling practices at the Gap, Inc. *Management Science* 68(11):7818–7836.
- Kremer M (1993) The O-ring theory of economic development. *Quarterly Journal of Economics* 108(3):551–575.
- Martinez-Moyano IJ, McCaffrey DP, Oliva R (2014) Drift and adjustment in organizational rule compliance. *Organization Science* 25(1):44–65.
- Mas A, Moretti E (2009) Peers at work. *American Economic Review* 99(1):112–145.
- McKnight DH, Choudhury V, Kacmar C (2002) Developing and validating trust measures for e-commerce: An integrative typology. *Information Systems Research* 13(3):334–359.
- Milgrom P, Roberts J (1990) The economics of modern manufacturing: Technology, strategy, and organization. *American Economic Review* 80(3):511–528.
- Mittal V, Ross WT, Baldasare PM (1998) The asymmetric impact of negative and positive attribute-level performance on overall satisfaction and repurchase intentions. *Journal of Marketing* 62(1):33–47.
- Oliva R, Sterman JD (2001) Cutting corners and working overtime: Quality erosion in the service industry. *Management Science* 47(7):894–914.
- Oliva R, Sterman JD, Giese M (2003) Limits to growth in the new economy: Exploring the ‘get big fast’ strategy in e-commerce. *System Dynamics Review* 19(2):83–117.

- Rahmandad H, Repenning N (2016) Capability erosion dynamics. *Strategic Management Journal* 37(4):649–672.
- Repenning NP, Sterman JD (2001) Nobody ever gets credit for fixing problems that never happened. *California Management Review* 43(4):64–88.
- Saghiri S, Wilding R, Mena C, Bourlakis M (2017) Toward a three-dimensional framework for omni-channel. *Journal of Business Research* 77:53–67.
- Sirdeshmukh D, Singh J, Sabol B (2002) Consumer trust, value, and loyalty in relational exchanges. *Journal of Marketing* 66(1):15–37.
- Sterman JD (2000) *Business Dynamics: Systems Thinking and Modeling for a Complex World*. Boston: Irwin/McGraw-Hill.
- Sun L, Abraham S (2021) Estimating dynamic treatment effects in event studies with heterogeneous treatment effects. *Journal of Econometrics* 225(2):175–199.
- Ton Z (2014) *The Good Jobs Strategy: How the Smartest Companies Invest in Employees to Lower Costs and Boost Profits*. New York: Houghton Mifflin Harcourt.
- Tversky A, Kahneman D (1992) Advances in prospect theory: Cumulative representation of uncertainty. *Journal of Risk and Uncertainty* 5(4):297–323.
- Tucker AL (2007) An empirical study of system improvement by frontline employees in hospital units. *Manufacturing & Service Operations Management* 9(4):492–505.