

# When Front-Stage Innovation Overloads Back-Stage Workers: A Dynamic Model of Workforce, Capability, and Customer Trust in Omnichannel Retail

(Authors' names are not included for peer review)

Authors are encouraged to submit new papers to INFORMS journals by means of a style file template, which includes the journal title. However, use of a template does not certify that the paper has been accepted for publication in the named journal. INFORMS journal templates are for the exclusive purpose of submitting to an INFORMS journal and are not intended to be a true representation of the article's final published form. Use of this template to distribute papers in print or online or to submit papers to another non-INFORM publication is prohibited.

**Abstract. Problem definition.** Digital investment in omnichannel retail generates back-stage demand on workers that standard dashboards cannot see. Front-stage automation amplifies operational complexity faster than back-stage capacity can absorb it, creating a structural measurement gap: technology returns and labour consequences appear on different ledgers, and the workforce that determines whether technology delivers its designed capability goes systematically underinvested.

**Methodology/results.** We formalise this gap as a parsimonious dynamic model and derive three analytically characterised results. (i) A computable digitalisation ceiling exists beyond which additional front-stage investment destroys customer lifetime value (CLV) by overloading back-stage workers; revenue metrics continue rising even as CLV collapses. (ii) The entire return to workforce investment flows through service capacity, invisible to dashboards that treat payroll as a cost line. (iii) The system is bistable: the neglected firm faces a safe rollout window less than two-fifths the size of its investing peer. We provide suggestive empirical support using panel data on 13 US and 11 European retailers (2010–2024): the technology–workforce interaction is positive in cross-section, and an inverted-U between e-commerce share and customer satisfaction aligns with the predicted digitalisation ceiling.

**Managerial implications.** The model supplies two computable numbers absent from standard dashboards: a digitalisation ceiling and a maximum safe rollout size. Both are workforce-contingent, making workforce investment a structural requirement of digital strategy rather than a social concession.

**Funding:** No external funding was received for this research.

**Key words:** omnichannel retail, workforce well-being, service capability, customer trust, system dynamics, people-centric operations

## 1. Introduction

Omnichannel grocery retailers have invested heavily in digital capabilities in recent years (Hübner et al. 2021, Bijmolt et al. 2021), enabling customers to build baskets across channels, reserve items for same-day collection, and track substitutions in real time. Yet as digital engagement has grown, retailers have encountered a structural asymmetry in how costs and returns are measured. Technology investment registers as successful in the front-stage domain where it is evaluated, while the back-stage workforce that must fulfil the digitally generated demand absorbs consequences that appear nowhere in the technology business case (Corbett 2024).

The asymmetry is structural, not incidental. Front-stage digital capabilities generate customer journeys that multiply handoff complexity far faster than order volume grows (Saghiri et al. 2017), placing distinctive operational demands on back-stage workers (pickers, drivers, click-and-collect associates). The empirical literature confirms the scale of these consequences: workforce investment generates returns that are both large and invisible to standard metrics (Fisher et al. 2021, Kesavan et al. 2022), and firms that invest in workforce capacity outperform on both operational and financial measures (Ton 2014), yet the causal chain runs through mechanisms that standard dashboards do not monitor.

What is missing is a formal account of the mechanism. The empirical literature establishes that workforce investment pays off; it does not characterise *when* additional front-stage investment begins to destroy rather than create value, *how fast* a firm can scale digitalisation before the back-stage workforce is overwhelmed, or *what dynamic relationship* connects technology investment, workforce well-being, and customer trust over time. These are not questions that regression analysis on historical data can answer: they require a model that treats workforce well-being and customer trust as co-evolving state variables responding to a controllable investment sequence (see §2).

This paper formalises the economic structure of this problem as a parsimonious dynamic model. Technology and workforce are not independent inputs to service capability but strategic complements (Milgrom and Roberts 1990, Bainbridge 1983, Kremer 1993). The central modelling innovation is the *workforce capacity multiplier*: effective service capability is the product of the technology stock and a workforce engagement index. When front-stage demand grows faster than back-stage workforce capacity, a service gap opens that depletes workforce well-being, which reduces effective capability, which widens the gap further: a reinforcing deterioration that single-domain models cannot generate.

Three formally characterised results follow. A *Front-Stage Trap* (Theorem 1): a computable digitalisation threshold exists beyond which additional front-stage investment reduces customer lifetime value by overloading back-stage workers. A *Good Jobs Buffer* (Theorem 2): the return to workforce investment flows through a capacity-level channel that is invisible to dashboards measuring only operational sensing. A *Burnout Cliff* (Theorem 3): the system exhibits bistability, and a firm that neglects its workforce faces a safe deployment window less than two-fifths the size of one that invests in its people. A fourth result, Proposition 4, establishes that the marginal value of workforce investment is increasing in digitalisation level. Workforce investment is therefore not a social concession but a structural requirement of digital strategy.

## 2. Related Literature

This paper sits at the intersection of three research streams. We review each and identify the gaps the model addresses.

*Service quality erosion.* The foundational insight that service quality can erode endogenously under demand pressure originates with Oliva and Serman (2001), whose system-dynamics model (Serman 2000) produces death spirals in which workers cut corners under demand pressure. Their finding that workers are roughly twice as willing to cut corners as to work overtime grounds the overload-asymmetry parameter  $\sigma$  in the present model. Subsequent work extends the logic to capability erosion (Rahmandad and Repenning 2016, Repenning and Serman 2001), e-commerce growth (Oliva et al. 2003), and regulatory compliance (Martinez-Moyano et al. 2014); KC and Terwiesch (2009) confirm that sustained high workload degrades service outcomes empirically, and Corbett (2024) calls for formal dynamic models of workforce well-being in operations. **Gap 1:** no model in this tradition treats workforce well-being as a positive capability multiplier. **Gap 2:** no Oliva–Serman extension adds customer trust, a front-/back-stage split, or digital scaling.

*Workforce investment and retail performance.* A substantial empirical literature documents the link between workforce practices and performance. The Good Jobs Strategy (Ton 2014) identifies four operational choices that drive outperformance; causal evidence comes from field experiments showing that responsible scheduling raises productivity and sales (Kesavan et al. 2022) and that profit-maximising staffing exceeds prevailing practice (Fisher et al. 2021, Dubé et al. 2020). The conceptual mechanism is the *service profit chain* (Heskett et al. 1994), which has not been formalised as a dynamic model with computable boundary conditions. The Job Demands-Resources model (Bakker and Demerouti 2007, Demerouti et al. 2001) provides the micro-foundation: overload depletes well-being faster than equivalent surplus restores it, and productivity responses

**Table 1 Literature gaps and model contributions.**

| Gap                                                                        | Closest prior work                                                  | This paper                                                                          |
|----------------------------------------------------------------------------|---------------------------------------------------------------------|-------------------------------------------------------------------------------------|
| 1. No model treats $W$ as a positive capability multiplier                 | Oliva and Sterman (2001): $W$ as degradation only                   | Effective capability as the product of technology and workforce engagement          |
| 2. No Oliva–Sterman extension adds trust, domain split, or digital scaling | Oliva and Sterman (2001); Martinez-Moyano et al. (2014)             | Customer trust as a co-evolving state; front/back domain split; digital scaling     |
| 3. Good Jobs Strategy / service profit chain (SPC) unformalised            | Ton (2014); Heskett et al. (1994): empirical only                   | Dynamic formalisation with analytical equilibria and boundary conditions            |
| 4. No cross-domain front-to-back overload model                            | Saghiri et al. (2017): conceptual only                              | Super-linear back-stage demand as a function of digital depth and customer trust    |
| 5. Digitalisation-workforce cross-partial not derived for services         | No direct service prior; see also Gans and Goldfarb (2026) (O-ring) | Formally derived: the return to workforce investment rises with the technology base |

scale with engagement (Mas and Moretti 2009). These two features (asymmetric overload and engagement-scaled sensitivity) underpin the workforce dynamics in §3.3. **Gap 3:** no dynamic model formalises the Good Jobs Strategy or the service profit chain with computable boundary conditions.

*Omnichannel capability.* Digitalisation creates considerable challenges for retail logistics, with fulfilment complexity scaling super-linearly as channel combinations increase (Hübner et al. 2021, Saghiri et al. 2017, Bell et al. 2018, Acimovic and Graves 2015). Automation shifts the operator's role toward exception handling (Bainbridge 1983), making the quality of residual human tasks more consequential for overall output (Gans and Goldfarb 2026). Technology–workforce strategic complementarity (Milgrom and Roberts 1990, Aral et al. 2012) motivates the model's central structural assumption that the marginal value of workforce engagement rises with the technology base (§3.1). **Gap 4:** no model formalises front-to-back cross-domain overload. **Gap 5:** the digitalisation-workforce cross-partial has not been derived for services.

Table 1 summarises the five gaps and the model's response to each.

### 3. Model

Sections 3.1–3.5 formalise the workforce capacity multiplier introduced in §1. The front-stage digitalisation level  $F > 0$  is treated as an exogenous managerial control: it sets the complexity and volume of customer-facing digital features, and thereby determines the demand placed on the back-stage operation. The model traces the dynamic consequences of choosing  $F$  for the three co-evolving back-stage state variables and, through them, for customer lifetime value (CLV). Each of

the eight free parameters has a single, distinct role with a clean kill condition (Table 2): setting it to its neutral value eliminates exactly one mechanism, allowing the model to nest single-domain benchmarks as special cases.

### 3.1. Stocks and Effective Capability

Three state variables evolve jointly:  $K(t) > 0$ , the back-stage technology stock (installed infrastructure, warehouse management systems, and process standardisation);  $W(t) \in (0, 1)$ , an index of workforce well-being (engagement, energy, and cognitive capacity among pickers, drivers, and click-and-collect associates); and  $T(t) \in (0, 1)$ , customer trust (the customer base's aggregate confidence that the retailer will deliver reliably). All three are dimensionless and normalised to their natural operating ranges. Throughout, starred variables ( $K^*$ ,  $W^*$ ,  $T^*$ ) denote steady-state equilibrium values; the pre-shock (baseline) digitalisation level is written  $F_0$ ; and  $\Delta F$  denotes an upgrade step, the increase in  $F$  from  $F_0$ .

The central structural innovation is the *effective capability*. A warehouse management system or substitution-handling process can only deliver its designed performance if the workforce operating it has the capacity to do so. Effective capability is therefore not the technology stock alone but its product with a workforce engagement multiplier:

$$\begin{aligned} K_{\text{eff}}(t) &= K(t) \cdot h(W(t)), \\ h(W) &= (1 - p) + pW, \quad p \in (0, 1). \end{aligned} \tag{1}$$

Effective capability is the product of the technology stock and a workforce capacity multiplier  $h(W)$ . The multiplier ranges from  $h(0) = 1 - p$  (at zero well-being, the workforce delivers only the irreducible, automation-driven fraction  $1 - p$  of the technology's designed potential) to  $h(1) = 1$  (a fully engaged workforce realises the technology's full design capacity). The cross-partial

$$\frac{\partial^2 K_{\text{eff}}}{\partial K \partial W} = p > 0 \tag{2}$$

is strictly positive for all  $p > 0$ , establishing technology and workforce as complements in the Milgrom and Roberts (1990) sense.

The parameter  $p$  governs complementarity:  $p = 0$  decouples workforce from capability;  $p = 1$  makes the operation fully labour-intensive. The baseline  $p = 0.5$ , motivated by capital-skill complementarity evidence (Krusell et al. 2000), represents a moderate starting point; estimation is discussed in §7.2. All three theorems hold qualitatively for any  $p \in (0, 1)$ .

### 3.2. Demand and the Service Gap

Customer demand on back-stage operations is

$$D(T, F) = (1 + aT) \cdot F^b, \quad (3)$$

where  $a > 0$  is the trust-demand coupling and  $b > 1$  is the front-stage amplification exponent. The trust-demand coupling reflects the established finding that customer trust drives spending consolidation and order frequency (Sirdeshmukh et al. 2002, McKnight et al. 2002). The super-linearity  $b > 1$  captures the well-documented increase in back-stage fulfilment complexity as front-stage channel breadth grows: as shown in §2, Bell et al. (2018) and Acimovic and Graves (2015) document patterns consistent with fulfilment error rates and demand-capacity mismatches rising faster than linearly with channel combinations. We set  $b = 1.3$  as an illustrative, structurally conservative baseline motivated by this literature; the Front-Stage Trap requires only  $b > 1$ , and all three theorems hold qualitatively across  $b \in [1.1, 2.0]$  (EC Table EC.3).

The *service gap* measures the shortfall between demand and effective capability:

$$G(t) = D(T(t), F) - K(t) \cdot h(W(t)). \quad (4)$$

When  $G > 0$  (overload), customers experience failures. When  $G < 0$  (surplus), service runs reliably. We write  $[G]^+ = \max(G, 0)$  and  $[-G]^+ = \max(-G, 0)$  for the positive and negative parts.

### 3.3. Dynamics

The feedback loops in §3.5 are built from the three ordinary differential equations (ODEs) introduced below, which translate each economic mechanism into a rate equation. All rate parameters carry units of months<sup>-1</sup>; each ODE is piecewise-smooth, activating different terms in the overload regime ( $G > 0$ , so  $[G]^+ = G$ ,  $[-G]^+ = 0$ ) and the surplus regime ( $G < 0$ , so  $[G]^+ = 0$ ,  $[-G]^+ = -G$ ).

*Capability.* The technology stock evolves as

$$\frac{dK}{dt} = r [G]^+ W - \delta K + u, \quad (5)$$

where  $r > 0$  is the sensing-adaptation rate,  $\delta > 0$  is the depreciation rate, and  $u \geq 0$  is direct back-stage investment. The sensing term  $r [G]^+ W$  captures organisational learning from service failures (Tucker 2007): the product structure requires both a gap to detect ( $[G]^+ > 0$ ) and workforce capacity to act on it ( $W > 0$ ), reflecting the finding that capability accumulation scales with the workforce knowledge stock (Argote and Epple 1990).

*Workforce well-being.* Well-being rises under surplus capacity and direct investment; it falls through natural decay and overload-amplified depletion:

$$\frac{dW}{dt} = ([-G]^+ + f)(1 - W) - (\mu + \sigma [G]^+)W. \quad (6)$$

Two grouped terms ensure  $W(t) \in [0, 1]$ . The first drives  $W$  upward (vanishing at  $W = 1$ ): capability surplus and direct workforce investment raise well-being with diminishing returns, consistent with the JD-R model (Bakker and Demerouti 2007). The second drives  $W$  downward (vanishing at  $W = 0$ ): natural decay ( $\mu$ ) and overload erosion ( $\sigma [G]^+$ ) deplete well-being, with  $\sigma > 1$  reflecting the asymmetry documented by Oliva and Sterman (2001). All three theorems hold for any  $\sigma > 1$  (EC Table EC.3).

*Customer trust.* Trust builds through reliable service and erodes faster when reliability fails (Sirdeshmukh et al. 2002, Mittal et al. 1998):

$$\frac{dT}{dt} = \beta [-G]^+ (1 - T) - \eta \beta [G]^+ T - \lambda T, \quad (7)$$

where  $\beta > 0$  is the trust-update rate,  $\eta > 1$  is the erosion asymmetry, and  $\lambda > 0$  is natural trust decay. We set  $\eta = 2.0$ , consistent with the asymmetric coefficients reported by Sirdeshmukh et al. (2002) (approximately 2.5:1 depletion-to-enhancement ratio).

The system (5)–(7) is piecewise-smooth; existence and uniqueness of solutions follow from Filippov's (1988) framework (details and smoothing robustness in EC §EC.1.1 and EC.1.8).

### 3.4. State-Space Invariance and Empirical Mapping

The state space  $K > 0$ ,  $W \in [0, 1]$ ,  $T \in [0, 1]$  is forward-invariant for all  $t \geq 0$ ; the verification via directional derivatives at each boundary is given in EC §EC.1.1.


Observable field proxies include: app feature index for front-stage capability ( $F$ ), employee engagement composite for workforce well-being ( $W$ ), net promoter score for customer trust ( $T$ ), and on-time fulfilment shortfall for the service gap ( $G$ ). Parameter calibration details are in EC §EC.4.

### 3.5. Feedback-Loop Architecture

Figure 1 shows the causal-loop architecture of the three ODEs, using the conventions of Sterman (2000). The diagram contains four feedback loops, each traceable from the state-variable arrows.

**Loop R1** ( $T \rightarrow D \rightarrow G \rightarrow T$ , regime-dependent: balancing in surplus, reinforcing under overload): trust raises demand, which changes the service gap, feeding back to trust. **Loop R2** ( $G \rightarrow W \rightarrow K_{\text{eff}} \rightarrow G$ , reinforcing): overload depletes well-being, reducing effective capability and

**Figure 1** Causal-loop architecture. Solid boxes are state variables; dashed boxes are exogenous controls. Arrow signs indicate direction of effect; loop polarity is shown in parentheses. R2 and R3, created by the capacity multiplier  $h(W)$ , are novel and do not exist in single-domain models.



widening the gap further — the deterioration loop behind the Front-Stage Trap. **Loop R3** ( $W \rightarrow K_{\text{eff}} \rightarrow G \rightarrow W$ , reinforcing): capability degradation escalates workforce depletion — the loop behind the Burnout Cliff. **Loop B1** ( $G \rightarrow K \rightarrow K_{\text{eff}} \rightarrow G$ , balancing): overload triggers sensing investment that raises capability and closes the gap. Loops R2 and R3, created by the capacity multiplier  $h(W)$ , have no counterpart in prior quality-erosion models.

### 3.6. Parameters and Special Cases

Table 2 summarises all parameters. Two control variables complete the model:  $u \geq 0$  (capability investment) and  $f \geq 0$  (workforce investment). The model recovers the core dynamics of Oliva and Sterman (2001) when  $p = 0$  and  $a = 0$  (see EC §EC.2 for the structural mapping), and reduces to a Solow-type capital-accumulation equation when  $p = 0$ ,  $a = 0$ , and  $r = 0$ .

## 4. Analytical Results

We derive four formal results: Theorems 1–3 characterise three phenomena that no single-domain model can produce; Proposition 4 establishes digitalisation-workforce complementarity. Full proofs are in the Electronic Companion (§EC.1); here we provide proof sketches.

### 4.1. Equilibrium Structure

The model generates two distinct long-run outcomes. The *surplus equilibrium*: effective capability exceeds demand, the workforce is healthy, trust builds, and the firm settles into a self-reinforcing high-performance state. The *low-trust attractor*: demand has overwhelmed capability, the workforce is depleted, and customers have withdrawn trust; exit requires structural intervention.

The technical characterisation of these two states is as follows. The surplus equilibrium exists for all  $F < F_{\text{crit}}$  and is unique under Condition C2 (EC Lemma 2, §EC.1.2). When effective capability exceeds demand, define the operational surplus  $S^* \equiv K^*h(W^*) - D(T^*, F) \geq 0$  as the excess of

**Table 2** Model parameters, baseline values, and kill conditions.

| # | Parameter                       | Baseline | Timescale                    | Kill                  | Mechanism killed     | Empirical anchor                                                   |
|---|---------------------------------|----------|------------------------------|-----------------------|----------------------|--------------------------------------------------------------------|
| 1 | $b$ (demand super-linearity)    | 1.3      | –                            | $b=1$                 | Front-Stage Trap     | Bell et al. (2018); Acimovic & Graves (2015)                       |
| 2 | $a$ (trust-demand coupling)     | 0.4      | –                            | $a=0$                 | Demand feedback      | McKnight et al. (2002); Sirdeshmukh et al. (2002)                  |
| 3 | $p$ (workforce complementarity) | 0.5      | –                            | $p=0$                 | Capacity multiplier  | Krusell et al. (2000)                                              |
| 4 | $\sigma$ (overload asymmetry)   | 2.0      | –                            | $\sigma=1$            | Tipping point        | Oliva and Sterman (2001); Bakker and Demerouti (2007) <sup>†</sup> |
| 5 | $\eta$ (erosion asymmetry)      | 2.0      | –                            | $\eta=1$              | Trust asymmetry      | Sirdeshmukh et al. (2002)                                          |
| 6 | $r$ (sensing-adaptation)        | 0.4      | $\tau_r \approx 1.7$ mo      | $r=0$                 | Self-correction dies | Tucker (2007); Argote & Epple (1990)                               |
| 7 | $\delta$ (depreciation)         | 0.03     | $\tau_\delta \approx 23$ mo  | –                     | –                    | Illustrative (EC §EC.4)                                            |
| 8 | $\beta$ (trust-update rate)     | 0.12     | $\tau_\beta \approx 5.8$ mo  | $\beta \rightarrow 0$ | Trust frozen         | Illustrative (EC §EC.4)                                            |
| + | $\mu$ (natural $W$ decay)       | 0.05     | $\tau_\mu \approx 14$ mo     | $\mu=0$               | $W$ exponential only | Demerouti et al. (2001); Dubé et al. (2020)                        |
| + | $\lambda$ (natural $T$ decay)   | 0.025    | $\tau_\lambda \approx 28$ mo | $\lambda=0$           | $T$ exponential only | Illustrative (EC §EC.4)                                            |

All time units are months. Timescales are  $\ln(2)/\theta$  (half-life). Eight free parameters (rows 1–8);  $\mu$  and  $\lambda$  govern natural decay. “Illustrative” parameters are calibrated within empirically informed ranges (EC §EC.4). <sup>†</sup>Magnitude comparator from professional services; see §3.3. Direct estimation of  $b$  and  $p$  from retailer panel data is the primary empirical extension (§7.2).

effective capability over demand at equilibrium. Equilibrium trust is an increasing function of  $S^*$ , saturating toward one as surplus grows:  $T^* = \beta S^* / (\beta S^* + \lambda)$ . Equilibrium well-being is similarly increasing in surplus and investment:  $W^* = (S^* + f) / (S^* + f + \mu)$ . The capability stock settles at  $K^* = u/\delta$ , determined entirely by the investment-depreciation balance. Full algebraic derivations are in EC §EC.1.2. Both  $T^*$  and  $W^*$  are increasing in  $S^*$ , so any improvement in effective capability propagates into higher trust and well-being at the equilibrium, creating the reinforcing dynamics that underpin the Good Jobs Buffer.

The low-trust attractor is reached whenever demand persistently overwhelms capability. With  $T^* \approx 0$ , demand collapses to its digitalisation-driven floor and the workforce settles at a low but strictly positive well-being level, held above zero by the workforce investment rate  $f$ . Effective capability at this attractor falls short of demand, maintaining a persistent residual gap. The attractor exists whenever  $F^b > (u/\delta) h(W_{\text{low}}^*)$  and is locally stable (EC Remark 1).

**Maintained conditions.** The formal results hold under four regularity conditions (C1)–(C4), verified analytically or numerically across the parameter space (EC Remark 3): surplus equilibrium existence, uniqueness, regularity, and low-trust attractor existence (EC §EC.1.2).

**REMARK 1 (ANALYTICAL TRANSPARENCY).** The table below classifies each result by evidentiary basis: Tier A (analytic for all valid parameters), Tier B (analytic under C1–C4), or Tier C (computational verification).

| Result                                                 | Tier | Basis                                                                                           |
|--------------------------------------------------------|------|-------------------------------------------------------------------------------------------------|
| Thm 1(i): equilibrium existence/nonexistence           | A    | Analytic (IVT + $\Phi(0)$ sign)                                                                 |
| Thm 1(ii): trajectories bounded                        | A    | Analytic (Gronwall; Lemma 3)                                                                    |
| Thm 1(ii): wide-basin convergence                      | C    | Numerical grid, $K_0 = K^*$ slice                                                               |
| Lemma 2 (uniqueness of $S^*$ )                         | B    | Analytic under C2-SC ( $\beta(f+\mu) < \lambda$ ; holds throughout parameter space; EC §EC.1.2) |
| Thm 2: $\partial \text{CLV}^* / \partial f > 0$ (sign) | B    | IFT chain; GE feedback correction sign analytic                                                 |
| Thm 2: GE feedback magnitude ( $\approx 0.09$ )        | C    | Confirmed across 100 configurations                                                             |
| Thm 3(i): $\Delta F_{\text{collapse}}$ threshold       | A    | Direct from Thm 1                                                                               |
| Thm 3(i): convergence from pre-shock IC                | C    | Numerical; basin-dependent                                                                      |
| Prop 4 cross-partial                                   | A    | Analytic ( $K_{\text{eff}} = K h(W)$ )                                                          |
| Prop 4 $F_{\text{max}}(W)$ boundary formula            | A    | Exact at $S^* \rightarrow 0$                                                                    |
| Prop 4 $F_{\text{max}}(W)$ interior upper bound        | B    | Analytic sign; $\approx 8\%$ correction numeric                                                 |
| Corollary 5, Case 1                                    | B    | Thm 1 + Thm 2 directly                                                                          |
| Corollary 5, Case 2 near $B_{\text{min}}$              | B    | Continuity from Case 1 (IFT; EC §EC.1.7)                                                        |
| Corollary 5, Case 2 interior                           | C    | Computational; $c_F/c_f \in [0.5, 5]$                                                           |

## 4.2. Theorem 1: The Front-Stage Trap

**THEOREM 1 (Front-Stage Trap).** *Suppose  $b > 1$  and  $p > 0$ . Define the digitalisation threshold*

$$F_{\text{crit}} = \left( \frac{u}{\delta} \cdot \left[ (1-p) + p \frac{f}{f+\mu} \right] \right)^{1/b}, \quad (8)$$


*which depends only on exogenous parameters  $(u, \delta, p, f, \mu, b)$  and equals  $K^* h(f/(f+\mu))^{1/b}$ , the capacity multiplier evaluated at the boundary workforce level. Then:*

*(i) For  $F < F_{\text{crit}}$ , the surplus equilibrium (unique under C2; Lemma 2, EC §EC.1.2) exists with  $T^* > 0$  and  $\text{CLV}^* > \text{CLV}_{\text{low}}$  (the CLV at the low-trust attractor).*

*(ii) For  $F \geq F_{\text{crit}}$ , no surplus equilibrium exists. All forward trajectories remain in a bounded compact set (Lemma 3, EC §EC.1.2), and the low-trust attractor is locally stable (EC Remark 1). The threshold  $F_{\text{crit}}$  is decreasing in  $p$  and increasing in  $f$ : the Trap binds earlier when workforce-technology coupling is strong, and is relaxed by workforce investment.*

*Proof sketch.* The surplus equilibrium requires  $\Phi(0) > 0$  (where  $\Phi(S) \equiv K^* h(W^*(S)) - D(T^*(S), F) - S$  is the surplus residual; full construction in EC §EC.1.2), which holds if and only if  $F < F_{\text{crit}}$ . At  $F = F_{\text{crit}}$ , the surplus  $S^* = 0$  and the equilibrium ceases to exist. For  $F > F_{\text{crit}}$ ,  $G > 0$ , activating overload depletion of  $W$  through Eq. (6), which reduces  $h(W)$ , which further widens  $G$  through Loop R2. The reinforcing dynamics drive the system toward the low-trust attractor. That trajectories remain bounded for all  $t$  follows from Lemma 3 (EC §EC.1.2): under overload,  $\dot{K} \leq C(F) - \delta K$  where  $C(F) = r(1+a)F^b + u$  is a demand-capacity bound, giving  $K(t) \leq \max(K_0, C(F)/\delta)$  by Gronwall's lemma; together with  $W, T \in [0, 1]$  by construction, the state space is forward-invariant and compact. Full proof in EC §EC.1.3.  $\square$

**Figure 2** Basin-of-attraction map. Green cells converge to the surplus attractor (high CLV) for  $F < F_{\text{crit}} \approx 1.11$ ; red cells converge to the low-trust attractor for  $F \geq F_{\text{crit}}$ . Doubled-grid validation confirms boundary stability (EC §EC.4.2).



At the calibrated baseline,  $F_{\text{crit}} \approx 1.11$ , corresponding to roughly a 40% increase in digital feature complexity above  $F_0 = 0.80$ . Figure 2 maps the regime structure: green cells converge to the surplus attractor (high CLV); red cells converge to the low-trust attractor (CLV  $\approx 400$ ). The cross-domain CLV sign reversal requires  $b > 1$  (super-linear demand) and  $p > 0$  (workforce-contingent threshold); at  $p = 0$  the threshold reduces to  $(u/\delta)^{1/b}$  and no longer responds to workforce investment.

#### 4.3. Theorem 2: The Good Jobs Buffer

**THEOREM 2 (Good Jobs Buffer).** *Suppose  $p > 0$  and the surplus equilibrium exists and is stable (Conditions C1–C4, i.e., the surplus equilibrium exists, is unique, and the low-trust attractor is present; Remark 1). Workforce investment  $f$  raises steady-state CLV through a multiplicative capacity-level channel:*

$$\frac{\partial \text{CLV}^*}{\partial f} > 0 \quad \text{for all } f \geq 0. \quad (9)$$

*The return decomposes into a direct capacity effect and an indirect trust effect:*

$$\begin{aligned} \frac{\partial \text{CLV}^*}{\partial f} = & \underbrace{\frac{\partial \text{CLV}}{\partial T^*} \cdot \frac{\partial T^*}{\partial S^*} \cdot K^* \cdot p}_{\text{direct: } h(W) \text{ raises } K_{\text{eff}}} \cdot \frac{\partial W^*}{\partial f} \\ & + \underbrace{\text{higher-order terms via sensing}}_{\text{indirect: } W \text{ in } r[G]^+ W}. \end{aligned}$$

*Proof sketch.* Since  $W^* = (S^* + f)/(S^* + f + \mu)$  at the surplus equilibrium, differentiating gives  $\partial W^*/\partial f = \mu/(S^* + f + \mu)^2 > 0$ : workforce investment raises steady-state well-being. Higher  $W^*$  raises  $h(W^*) = (1 - p) + pW^*$ , which increases  $K_{\text{eff}} = K^* \cdot h(W^*)$ , which widens the surplus  $S^*$ , which raises  $T^* = \beta S^*/(\beta S^* + \lambda)$ , which raises CLV. Each link is strictly positive:  $\partial h/\partial W = p > 0$ ,  $\partial S/\partial h = K^* > 0$ ,  $\partial T^*/\partial S^* > 0$ ,  $\partial \text{CLV}/\partial T^* > 0$ . The chain-rule product is therefore positive, and diminishing returns follow from the concavity of  $W^*$  in  $f$ . Full proof in EC §EC.1.4.  $\square$

The  $p=0$  counterfactual in §5.1 confirms this: disabling the workforce multiplier leaves equilibrium CLV exactly invariant to workforce spending. The entire CLV return flows through service capacity, not the operational sensing channel. Standard dashboards, which record technology and workforce returns in separate ledgers, structurally underestimate the financial return to workforce investment.

#### 4.4. Theorem 3: The Burnout Cliff

**THEOREM 3 (Burnout Cliff).** *Suppose  $\sigma > 1$ ,  $p > 0$ , and  $F_0 < F_{\text{crit}}$  (a surplus equilibrium exists at  $F_0$ ). Under these conditions the system is bistable in a neighbourhood of  $F_{\text{crit}}$ . Consider a firm at the surplus equilibrium that implements a step increase in  $F$  from  $F_0$  to  $F_0 + \Delta F$ . Two thresholds govern the outcome:*

(i) *Permanent-collapse onset. For  $\Delta F \geq \Delta F_{\text{collapse}}$ , no surplus equilibrium exists at the post-upgrade level. The low-trust attractor is the only locally stable state (EC Remark 1); convergence to it is confirmed numerically across the studied initial conditions.*

$$\begin{aligned} \Delta F_{\text{collapse}} &= F_{\text{crit}} - F_0 \\ &= \left( \frac{u}{\delta} [(1-p) + p \frac{f}{f+\mu}] \right)^{1/b} - F_0. \end{aligned}$$

For  $\Delta F < \Delta F_{\text{collapse}}$ , a new surplus equilibrium exists at the post-upgrade level. Whether a given trajectory reaches this equilibrium depends on its initial condition relative to the basin boundary (Figure 2); numerical verification confirms convergence from the pre-shock equilibrium initial condition across the parameter configurations studied (EC §EC.4.2).

(ii) Instantaneous-overload diagnostic. A second, larger threshold marks the boundary above which demand turns positive at  $t = 0^+$ :

$$\Delta F_{\text{rapid}} = \frac{S_0}{b \cdot F_0^{b-1} \cdot (1 + aT^*)}, \quad (10)$$

where  $S_0 = K^*h(W^*) - D(T^*, F_0)$  is the pre-upgrade surplus. Because  $\Delta F_{\text{rapid}} > \Delta F_{\text{collapse}}$  (at baseline: 0.40 vs. 0.31), a firm in the zone  $\Delta F \in [\Delta F_{\text{collapse}}, \Delta F_{\text{rapid}})$  collapses permanently but through a slow trajectory ( $G$  negative at  $t = 0^+$ , surplus eroded gradually). For  $\Delta F \geq \Delta F_{\text{rapid}}$ ,  $G$  turns positive immediately and the system enters a rapid self-reinforcing collapse to the low-trust attractor.

The safe zone is  $\Delta F < \Delta F_{\text{collapse}}$ ;  $\Delta F_{\text{rapid}}$  is a deployment-speed diagnostic, not a safety threshold.

*Proof sketch.* Part (i) follows directly from Theorem 1: for  $F_0 + \Delta F \geq F_{\text{crit}}$ , no surplus equilibrium exists; for  $\Delta F < \Delta F_{\text{collapse}}$ , a new surplus equilibrium exists at  $F_0 + \Delta F$ , and numerical verification confirms that trajectories originating at the pre-shock surplus equilibrium converge to this new surplus equilibrium across the parameter configurations studied (EC §EC.4.2). For Part (ii), setting the first-order demand increment  $\Delta D \approx (1 + aT^*) b F_0^{b-1} \Delta F$  equal to the pre-upgrade surplus  $S_0$  gives  $\Delta F_{\text{rapid}}$ . Under overload the workforce tipping threshold (the  $W$  level below which depletion rate exceeds recovery rate) is  $W_{\text{tip}} = f/(f + \mu + \sigma G_0^+)$ ; when  $\Delta F < \Delta F_{\text{rapid}}$ ,  $G_0^+ = 0$  at  $t_0^+$  and the system crosses  $W_{\text{tip}}$  through a slow transient, consistent with the slow-collapse zone. Full proof in EC §EC.1.5.  $\square$

Workforce investment  $f$  raises  $F_{\text{crit}}$  and enlarges the safe upgrade zone. Before each deployment phase, the firm should compute  $\Delta F_{\text{collapse}} = F_{\text{crit}}(f) - F_0$  as a hard ceiling on planned upgrade size; this ceiling is set entirely by  $f$ .

#### 4.5. Proposition 4: Digitalisation-Workforce Complementarity

**PROPOSITION 4 (Digitalisation-Workforce Complementarity).** *The marginal value of workforce improvement is strictly increasing in the technology base:*

$$\begin{aligned}\frac{\partial K_{\text{eff}}}{\partial W} &= p \cdot K, \\ \frac{\partial^2 K_{\text{eff}}}{\partial W \partial K} &= p > 0.\end{aligned}\tag{11}$$

*Since maintaining surplus at higher  $F$  requires higher effective capability  $K_{\text{eff}}$  (that is, either higher  $K^*$  through increased  $u$ , or higher  $h(W^*)$  through increased  $f$ ), the equilibrium return to workforce investment increases with digitalisation level.*

*Equivalently, workforce well-being determines the maximum sustainable digitalisation level. At the regime boundary ( $S^* \rightarrow 0$ ), the equilibrium trust approaches zero ( $T^* \rightarrow 0$ ; see Lemma 2, EC §EC.1.2), so:*

$$F_{\text{max}}(W) = (K^* \cdot h(W))^{1/b},\tag{12}$$

*which is strictly increasing in  $W$  for  $p > 0$ . The full implicit differentiation accounting for  $T^*(W, F)$  dependence at interior equilibria (where  $S^* > 0$ ) yields a correction term of order  $O(aS^*)$  that vanishes as  $S^* \rightarrow 0$ ; the formula (12) is therefore exact at the boundary and an upper bound on  $F_{\text{max}}$  in the interior: trust-elevated demand means true safe capacity is somewhat lower than this formula implies (EC §EC.1.6).*


*Proof.* The cross-partial follows directly from  $K_{\text{eff}} = K[(1-p) + pW]$ . The  $F_{\text{max}}$  formula identifies  $F_{\text{crit}}$  at  $S^* \rightarrow 0$ , where  $T^* \rightarrow 0$ , giving  $(K^*h(W))^{1/b}$ , strictly increasing in  $W$  since  $h'(W) = p > 0$ ; with the full  $T^*(W, F)$  correction this is exact at the boundary and an upper bound in the interior (EC §EC.1.6).  $\square$

Figure 3 illustrates: the gap between a stressed firm ( $W = 0.30$ ,  $F_{\text{max}} \approx 1.063$ ) and a healthy one ( $W = 0.70$ ,  $F_{\text{max}} \approx 1.307$ ) is 23%, arising entirely from workforce investment.

**COROLLARY 5 (Budget-Constrained Trap).** *Under a fixed budget constraint  $c_F F + c_f f \leq B$ , the CLV-maximising allocation always includes positive workforce investment  $f^* > 0$ : the all- $F$  corner is strictly dominated whenever the budget reaches or exceeds the Trap threshold.*

*Proof.* See EC §EC.1.7.  $\square$

**Figure 3** Digitalisation-workforce complementarity. The healthy firm ( $W = 0.70$ ) sustains 23% higher  $F_{\max}$  than the stressed firm ( $W = 0.30$ ), arising entirely from workforce investment.



## 5. Computational Study


We illustrate the analytical results through calibrated simulations using the baseline parameters in Table 2. CLV follows Gupta et al. (2004):  $CLV = w(T)/(1 - \rho(T))$ , where  $\rho(T)$  is an annual retention probability and  $w(T)$  an annualised per-customer margin, with  $\rho(T) = 0.80 + 0.18T$  and  $w(T) = 80(1 + 0.25 \max(T - 0.15, 0))$ . The endpoints  $\rho(0) = 0.80$  and  $\rho(1) = 0.98$  fall within the 70–90% annual range in Gupta et al. (2004); £80 is an illustrative annual margin for UK omnichannel grocery. All analytical results hold for any CLV increasing in trust (calibration in EC §EC.4.3). All three qualitative phenomena are robust across  $b \in [1.1, 2.0]$ ,  $\sigma \in [1.0, 3.0]$ , and  $p \in [0.25, 0.75]$  (EC Table EC.3).

### 5.1. The Hidden Cost of App Innovation (Figure 4)

Three deployment strategies span the critical threshold  $F_{\text{crit}} \approx 1.11$ . All begin from  $(K_0, W_0, T_0) = (1.60, 0.55, 0.50)$  with  $F_0 = 0.80$ . Scenario A ramps aggressively to  $F_{\text{end}} = 1.15$  with no workforce investment; Scenario B ramps to  $F_{\text{end}} = 1.05$  with concurrent support ( $f = 0.06$ ); Scenario C pre-invests in workforce ( $f = 0.10$ , months 1–12, after which  $f$  returns to 0.0) before ramping to  $F_{\text{end}} = 1.15$ .

All three appear profitable in front-stage KPIs (panel a), yet CLV diverges sharply (panel b): by month 24, Scenario C sustains CLV approximately 40% above Scenario A. Workforce collapse

**Figure 4** The hidden cost of app innovation. Panel (a): the demand-volume index  $F \cdot D(T(t), F(t))$  rises for all three scenarios, masking divergent back-stage outcomes. Panel (b): CLV diverges sharply by month 24. Panels (c,d): workforce collapse precedes trust collapse in the no-investment scenario.



precedes trust collapse (panels c, d), confirming the cross-domain cost invisible in the technology business case.

## 5.2. Basin-of-Attraction Structure


Figure 2 maps the basin boundary across the joint space of  $W_0$  and  $F$ . The boundary shifts rightward as  $W_0$  increases: a firm near the basin boundary can increase its safety margin by raising workforce investment  $f$ , which shifts  $F_{\text{crit}}$  rightward without touching the technology stock.

## 5.3. The Burnout Cliff (Figure 5)

We sweep  $\Delta F \in [0, 0.60]$  for three workforce-investment scenarios: neglected ( $f = 0.01$ ), baseline ( $f = 0.03$ ), and invested ( $f = 0.10$ ), each starting from its surplus equilibrium at  $F_0 = 0.80$ . Varying  $f$  shifts  $F_{\text{crit}}(f)$  from 0.98 to 1.29, widening  $\Delta F_{\text{collapse}}$  from 0.18 to 0.49.

Figure 5 shows the three normalised CLV curves. For the neglected firm ( $f = 0.01$ ), the safe zone collapses to  $\Delta F < 0.18$ ; for the invested firm ( $f = 0.10$ ), it extends to  $\Delta F = 0.49$ . The slow-

**Figure 5 The Burnout Cliff. Normalised terminal CLV versus upgrade size  $\Delta F$  for three workforce-investment levels. Dotted verticals:  $\Delta F_{\text{collapse}}$  (0.18, 0.31, 0.49). The neglected firm ( $f = 0.01$ ) exhausts its safe zone at  $\Delta F = 0.18$ ; the invested firm ( $f = 0.10$ ) retains headroom to  $\Delta F = 0.49$ .**



collapse zone between  $\Delta F_{\text{collapse}}$  and  $\Delta F_{\text{rapid}}$  is the diagnostic blind spot: dashboards report continued growth while the trust trajectory has already turned negative. For the well-investing firm, this zone disappears entirely (EC §EC.1.5).

## 6. Empirical Evidence

The model yields three testable predictions at the firm–year level: (i) a positive technology–workforce interaction ( $K \times W > 0$ ); (ii) an inverted-U between e-commerce intensity and customer satisfaction; (iii) a positive effect of workforce improvement on satisfaction that increases with technology intensity. We explore patterns consistent with each prediction using panel data from two independent geographies.

### 6.1. Data and Sample

*US sample.* We construct a balanced panel of 13 publicly listed US retailers observed annually from 2010 to 2024. Customer satisfaction ( $T$  proxy) is measured by the American Customer Satisfaction Index (ACSI; Fornell et al. 1996), a nationally representative survey scored 0–100 (source: [theacsi.org](http://theacsi.org)). ACSI is the most widely used satisfaction metric in operations management (Fisher et al. 2021, Kesavan et al. 2022) and has established predictive validity for future financial performance. Workforce well-being ( $W$  proxy) uses Glassdoor overall ratings (1–5 scale), an

imperfect but increasingly standard proxy for aggregate employee sentiment in the absence of proprietary engagement surveys (Corbett 2024). Technology investment ( $K$  proxy) is capital expenditure intensity (CapEx/Revenue) extracted from SEC EDGAR 10-K filings; this captures total capital investment rather than digital-specific spending, introducing measurement noise that biases interaction coefficients toward zero. E-commerce share of retail sales is from the Federal Reserve Economic Data series ECOMPCTSA. Revenue (size control), SGA intensity, and net margin are from EDGAR. After listwise deletion of missing observations, the regression sample contains 131 complete firm–year observations across 13 firms and 12 years.

*EU sample.* An independent replication panel covers 11 European retailers across six countries (UK, France, Germany, Spain, Sweden, Netherlands) from 2019 to 2023. Customer satisfaction uses Trustpilot scores (1–5 scale; rescaled to 0–100 for comparability). Glassdoor ratings proxy workforce well-being. Financial data are from Yahoo Finance. Because financials are reported in mixed currencies (GBP, EUR, SEK), we use within-year asset percentile rank as the size control in cross-sectional specifications and log revenue only in within-firm (fixed-effects) specifications where currency is absorbed. E-commerce share uses UK Office for National Statistics internet sales data for UK firms and Eurostat data for continental retailers. The complete panel contains 20 firm–year observations with all three proxies present.

## 6.2. Identification Strategy

All variables are standardised to mean zero and unit standard deviation before regression, following the recommendation of Aral et al. (2012) for interaction terms in small samples. Standardisation centres the interaction at the sample mean, improving interpretability and reducing multicollinearity between the interaction and its constituent terms. We estimate three specifications mapping to the three model predictions:

*Test 1 (Complementarity  $\rightarrow$  Proposition 4).*

$$T_{it} = \alpha_i + \gamma_t + \beta_1 K_{it}^z + \beta_2 W_{it}^z \quad (13)$$

$$+ \beta_3 (K \times W)_{it}^z + \mathbf{X}_{it} \boldsymbol{\delta} + \varepsilon_{it}. \quad (14)$$

Prediction:  $\beta_3 > 0$ . Two-way fixed effects (firm  $\alpha_i$  and year  $\gamma_t$ ) absorb time-invariant firm heterogeneity and common shocks. Controls  $\mathbf{X}_{it}$  include log revenue, and in extended specifications, SGA intensity and net margin.

*Test 2 (Front-Stage Trap → Theorem 1).*

$$T_{it} = \alpha_i + \beta_1 \text{Ecom}_{it} + \beta_2 \text{Ecom}_{it}^2 + \mathbf{X}_{it} \boldsymbol{\delta} + \varepsilon_{it}. \quad (15)$$

Prediction:  $\beta_1 > 0$ ,  $\beta_2 < 0$  (inverted-U). Turning point:  $\text{Ecom}^* = -\beta_1/(2\beta_2)$ .

*Test 3 (Good Jobs Buffer → Theorem 2).*

$$\Delta T_{it} = \alpha_i + \gamma_t + \beta_1 \Delta W_{it} \quad (16)$$

$$+ \beta_2 (\Delta W \times K)_{it} + \mathbf{X}_{it} \boldsymbol{\delta} + \varepsilon_{it}. \quad (17)$$

Prediction:  $\beta_2 > 0$  (workforce improvement matters more at high technology intensity).

*Inference corrections.* With only  $N_c = 13$  (US) and  $N_c = 11$  (EU) firm clusters, standard clustered standard errors are biased downward (Cameron et al. 2008). We adopt three progressively conservative corrections. First, all reported  $p$ -values use the  $t(N_c - 1)$  distribution rather than the standard normal, which penalises small cluster counts. Second, we supplement with CR2 finite-sample bias-corrected standard errors (Bell and McCaffrey 2002), which inflate variance estimates to compensate for leverage effects in unbalanced panels. Third, we report Webb six-point wild cluster bootstrap  $p$ -values (Webb 2023) in the EC; this distribution is specifically designed for settings with  $N_c < 20$  clusters where standard wild bootstrap weight distributions (Rademacher, Mammen) produce unreliable rejection frequencies. All corrections are applied uniformly across both geographies. The most conservative inference (Webb bootstrap) is used as the primary benchmark for significance claims throughout.

### 6.3. Results

Table 3 reports the US results.

*Test 1 (Complementarity).* The cross-sectional (OLS) estimate of  $\beta_3$  is positive (+1.95), consistent with the model's prediction that firms with jointly higher technology and workforce investment enjoy higher satisfaction. This finding is economically meaningful: a one-standard-deviation joint increase in  $K$  and  $W$  is associated with a 1.95-point ACSI gain, comparable in magnitude to the main effects of  $K$  and  $W$  individually. The direction aligns with the complementarity evidence in Aral et al. (2012), who report similar positive cross-terms for technology–human-resource interactions in a different (non-retail) setting.

**Table 3 US regression results. Dependent variable: ACSI score (0–100). All  $p$ -values use  $t(N_c - 1)$  correction with  $N_c = 13$ –17 clusters. Significance: \*\*\* $p < 0.01$ , \*\* $p < 0.05$ , \* $p < 0.10$ .**

|                                 | 1a: OLS             | 1b: Firm FE             | 1c: Two-way FE | 1d: FE + controls |
|---------------------------------|---------------------|-------------------------|----------------|-------------------|
| <i>Test 1: Complementarity</i>  |                     |                         |                |                   |
| $K^z$                           | 0.2903              | −0.2153                 | −0.4218        | −0.3102           |
| $W^z$                           | 1.2506              | 0.1087                  | 0.3181         | 0.4553            |
| $K \times W^z$                  | 1.9533              | −0.6342**               | −0.5615*       | −0.7649           |
| $N$                             | 131                 | 131                     | 131            | 83                |
| <i>Test 2: Front-Stage Trap</i> |                     |                         |                |                   |
|                                 | 2a: Linear          | 2b: Quadratic           | 2c: Quad+W     |                   |
| Ecom                            | 0.4379              | 0.4400                  | 0.0284         |                   |
| Ecom <sup>2</sup>               |                     | −0.0279**               | −0.0055        |                   |
| Turning point                   |                     | 7.9%                    | 2.6%           |                   |
| $N$                             | 216                 | 216                     | 190            |                   |
| <i>Test 3: Good Jobs Buffer</i> |                     |                         |                |                   |
|                                 | 3a: $\Delta W$ only | 3b: $\Delta W \times K$ | 3c: + controls |                   |
| $\Delta W \times K$             |                     | −41.40                  | −132.37        |                   |
| $N$                             | 173                 | 118                     | 84             |                   |

$N$  varies across tests because each requires different variable coverage: Test 1 requires ACSI, Glassdoor, and CapEx; Test 2 requires only ACSI and e-commerce share; Test 3 uses first-differenced variables, losing one observation per firm.

However, the within-firm estimates (Models 1b–1d) are negative and sometimes significant, suggesting that year-to-year co-movements in  $K$  and  $W$  within the same firm do not raise satisfaction. This reversal is consistent with low within-firm variation in Glassdoor ratings (within-firm SD is only 30% of total SD): firm fixed effects absorb the cross-sectional signal that drives complementarity. A between-effects specification (Bell et al. 2018) would be more appropriate for identifying a structural relationship that varies primarily across firms, but the small number of clusters ( $N_c = 13$ ) limits the reliability of such estimates. The sign reversal under fixed effects is a common feature of complementarity tests when the identifying variation is predominantly cross-sectional (Milgrom and Roberts 1990).

*Test 2 (Front-Stage Trap).* The quadratic specification (2b) yields  $\hat{\beta}_2 = -0.028$  ( $p = 0.028$  under  $t(16)$ ), confirming the predicted inverted-U. The estimated turning point of 7.9% e-commerce share corresponds to the early-to-mid 2010s transition, broadly consistent with the model's  $F_{\text{crit}} \approx 1.11$  threshold. The negative quadratic term is robust to the inclusion of workforce controls (Model 2c), though the coefficient attenuates and loses significance ( $p = 0.31$ ), consistent with the model's prediction that the inverted-U is partially mediated by workforce overload. This attenuation is expected: the model predicts that e-commerce intensity damages satisfaction *through* workforce depletion, so controlling for  $W$  should absorb part of the quadratic effect. The turning point itself (7.9%) is

plausible given that US e-commerce share crossed this threshold around 2016 (FRED series ECOM-PCTSA), a period when several major retailers reported significant click-and-collect growing pains (Hübner et al. 2021).

*Test 3 (Good Jobs Buffer).* The interaction  $\Delta W \times K$  is negative in both specifications, contrary to the prediction. The first-differenced design, combined with noisy annual Glassdoor ratings and few complete observations ( $N = 84\text{--}118$ ), leaves this test underpowered. The Job Demands-Resources literature (Bakker and Demerouti 2007) suggests that well-being changes operate on quarterly or monthly timescales, whereas annual Glassdoor ratings capture a lagged, smoothed signal. A store-level test with higher-frequency engagement data would provide a more appropriate test of the capacity-multiplier channel.

*EU replication.* An independent European sample ( $N_c = 11$  retailers, 20 firm-year observations) yields a directionally consistent but statistically insignificant complementarity interaction.<sup>1</sup> The qualitative sign consistency across geographies—positive  $K \times W$  in cross-section—provides suggestive, though not conclusive, support for complementarity.

#### 6.4. Assessment

The three tests produce a mixed but informative pattern: the inverted-U (Test 2) is the strongest finding, the cross-sectional complementarity (Test 1) is directionally consistent, and the within-firm dynamics test (Test 3) is underpowered. We interpret the evidence as suggestive support for the model's qualitative predictions, subject to the limitations discussed in §7.2.

### 7. Discussion

#### 7.1. Contributions and Managerial Implications

The three theorems together provide the formal micro-foundation that connects Ton's (2014) four operational choices to each other through the dynamics of  $K_{\text{eff}}$ . Her prescription to “operate with slack” corresponds to the positive workforce investment rate  $f$ : operational slack keeps  $W$  from depleting under moderate overload, and Theorem 2 quantifies the CLV return to doing so. “Standardise and empower” corresponds to a lower overload sensitivity  $\sigma$ : well-designed processes reduce the per-unit depletion rate and enlarge the safe upgrade zone of Theorem 3. “Offer less” maps directly onto keeping  $F$  below  $F_{\text{crit}}$ : the model gives this demand-side lever a computable threshold

<sup>1</sup> Using Trustpilot (1–5 scale) as the outcome, the between-effects estimate of  $\beta_3$  is positive (+0.47) and significant under pooled cluster-robust inference (CR2:  $p = 0.045$  at  $t(10)$ ). However, the conservative HC3 between-effects estimate ( $p = 0.53$ ) and Webb bootstrap ( $p = 0.62$ ) do not reject zero, reflecting the extreme power limitation of  $N_c = 11$  clusters (Cameron et al. 2008). Full details in EC §EC.5.

rather than a qualitative prescription. “Cross-train” maps onto a higher complementarity parameter  $p$ : cross-trained workers raise effective capability  $h(W)$  at any given engagement level, increasing the marginal return to both technology and workforce investment simultaneously (Proposition 4). The model extends Oliva–Sterman with customer trust, a front-/back-stage domain split, and digital scaling: three extensions absent from the quality-erosion framework, now formalised in a single model that, to the best of our knowledge, provides the first dynamic formalisation of the service profit chain (Heskett et al. 1994) and Good Jobs Strategy (Ton 2014) with computable boundary conditions.

**Managerial implications.** Three implications follow. First, the maximum safe rollout speed is a workforce constraint (Theorem 3): the firm should verify that  $\Delta F_{\text{collapse}} = F_{\text{crit}}(f) - F_0$  exceeds the planned upgrade size. Second, workforce well-being is a leading indicator:  $W$  reaches crisis levels 6–21 months before trust does (EC Table EC.4). Third, the more digital the firm becomes, the more it should invest in its people (Proposition 4): conventional ROI calculations underestimate returns because the dominant capacity-multiplier channel is invisible in standard dashboards (Theorem 2).

## 7.2. Limitations and Directions for Future Research

*Model limitations.* The model treats  $F$  as exogenous; endogenising it would enable optimal deployment sequencing. The omission of a direct  $W \rightarrow D$  channel (whereby disengaged workers reduce throughput directly, independent of the service gap) makes our results a conservative lower bound on true complementarity. The linear form  $h(W) = (1 - p) + pW$  is a tractability choice; a concave specification (e.g.,  $h(W) = (1 - p) + pW^\gamma$ ,  $\gamma < 1$ ) would strengthen the Trap by making capability more sensitive to well-being at low  $W$  levels.

*Empirical limitations.* The firm-level panel analysis (§6) provides suggestive but not definitive support. Five specific concerns warrant emphasis. First, with only  $N_c \in \{11, 13\}$  firm clusters, clustered inference is at the boundary of asymptotic validity (Cameron et al. 2008); we report Webb six-point wild cluster bootstrap  $p$ -values in the EC for all key coefficients (Webb 2023). Second, Glassdoor ratings exhibit low within-firm temporal variation (within-firm SD is only 30% of total SD), making the complementarity interaction difficult to identify within firm over time; between-effects estimates that exploit cross-sectional variation are more informative but face omitted-variable concerns. Third, the EU outcome (Trustpilot) is complaint-biased and operates on a different scale than ACSI (Fornell et al. 1996); while we rescale to 0–100 for coefficient comparability, measurement equivalence is not guaranteed. Fourth, the complementarity interaction reverses sign under firm

fixed effects (Models 1b–1d), consistent with limited within-firm variation rather than absence of the effect; store-level panel data with exogenous technology rollout events would permit sharper separation of the capacity-multiplier channel from confounds. Fifth, the data are at the firm–year level; the model’s predictions concern service-capacity dynamics that are most cleanly testable with store-level panel data, high-frequency workforce surveys, and exogenous technology rollout events (e.g., staggered click-and-collect deployment across stores of the same chain).

*Extensions.* Three, in order of expected impact: (i) store-level identification of  $p$  using technology rollout timing as instruments; (ii) a competitive multi-firm extension where workforce poaching creates strategic complementarity; (iii) stochastic demand forcing, which would transform the Burnout Cliff into a probabilistic threshold.

## 8. Conclusion

Standard operational dashboards measure the front-stage domain that invests in technology, not the back-stage workforce that fulfils the resulting demand and absorbs the consequences. This disconnect is not merely a measurement problem; it is a structural feature of omnichannel operations that makes cross-domain value destruction structurally predictable from two observable managerial choices: the workforce investment rate  $f$  and the digitalisation level  $F$ . This paper provides the formal instrument for both.

Three results follow that, to our knowledge, no existing single formulation produces. The Front-Stage Trap (Theorem 1): a computable digitalisation threshold  $F_{\text{crit}} \approx 1.11$  beyond which additional front-stage investment destroys CLV by overloading back-stage workers, a sign reversal impossible in single-domain models. The Good Jobs Buffer (Theorem 2): workforce investment raises CLV through a capacity-level channel confirmed exact by the  $p=0$  counterfactual, which leaves CLV invariant to workforce spending and isolates the entire gain to service capacity. The Burnout Cliff (Theorem 3): the maximum safe upgrade size ranges from  $\Delta F = 0.18$  for a workforce-neglecting firm to  $\Delta F = 0.49$  for an investing firm, an approximately 172% difference; the neglected firm retains approximately 37% of the safe deployment headroom.

These results give operations teams two computable numbers standard dashboards do not report: a digitalisation ceiling ( $F_{\text{crit}}$ ) and a maximum safe rollout size ( $\Delta F_{\text{collapse}}$ ). Firm-level panel evidence from 13 US and 11 EU retailers provides suggestive support: the technology–workforce interaction is positive in cross-section (both geographies), and an inverted-U between e-commerce share and customer satisfaction aligns with the predicted digitalisation ceiling. Store-level identification

of  $p$  using exogenous technology rollout events remains the most important empirical follow-up. The broader implication extends beyond omnichannel retail: wherever automated front-stage processes generate demand on human back-stage operations, workforce investment is a structural requirement, not a social concession.

## References

- Acimovic J, Graves SC (2015) Making better fulfillment decisions on the fly in an online retail environment. *Manufacturing & Service Operations Management* 17(1):34–51.
- Aral S, Brynjolfsson E, Wu L (2012) Three-way complementarities: Performance pay, human resource analytics, and information technology. *Management Science* 58(5):913–931.
- Argote L, Epple D (1990) Learning curves in manufacturing. *Science* 247(4945):920–924.
- Bainbridge L (1983) Ironies of automation. *Automatica* 19(6):775–779.
- Bakker AB, Demerouti E (2007) The Job Demands-Resources model: State of the art. *Journal of Managerial Psychology* 22(3):309–328.
- Bell RM, McCaffrey DF (2002) Bias reduction in standard errors for linear regression with multi-stage samples. *Survey Methodology* 28(2):169–181.
- Bell DR, Gallino S, Moreno A (2018) Offline showrooms in omnichannel retail: Demand and operational benefits. *Management Science* 64(4):1629–1651.
- Bijmolt THA, Broekhuis M, de Leeuw S, Hirche C, Roederkerk RP, Sousa R, Zhu SX (2021) Challenges at the marketing-operations interface in omni-channel retail environments. *Journal of Business Research* 122:864–874.
- Cameron AC, Gelbach JB, Miller DL (2008) Bootstrap-based improvements for inference with clustered errors. *Review of Economics and Statistics* 90(3):414–427.
- Corbett CJ (2024) OM Forum: The operations of well-being: An operational take on happiness, equity, and sustainability. *Manufacturing & Service Operations Management* 26(2):409–430.
- Demerouti E, Bakker AB, Nachreiner F, Schaufeli WB (2001) The Job Demands-Resources model of burnout. *Journal of Applied Psychology* 86(3):499–512.
- Dubé A, Reich A, Shierholz H, West M (2020) Minimum wage shocks, employment flows and labor market frictions. *Journal of Labor Economics* 38(4):1009–1052.
- Filippov AF (1988) *Differential Equations with Discontinuous Righthand Sides*. Dordrecht: Kluwer Academic Publishers.
- Fisher ML, Gallino S, Netessine S (2021) Setting retail staffing levels: A methodology validated with implementation. *Manufacturing & Service Operations Management* 23(6):1562–1579.
- Fornell C, Johnson MD, Anderson EW, Cha J, Bryant BE (1996) The American Customer Satisfaction Index: Nature, purpose, and findings. *Journal of Marketing* 60(4):7–18.

- Gans JS, Goldfarb A (2026) O-ring automation. NBER Working Paper No. 34639. National Bureau of Economic Research, Cambridge, MA.
- Gupta S, Lehmann DR, Stuart JA (2004) Valuing customers. *Journal of Marketing Research* 41(1):7–18.
- Heskett JL, Jones TO, Loveman GW, Sasser Jr WE, Schlesinger LA (1994) Putting the service-profit chain to work. *Harvard Business Review* 72(2):164–174.
- Hübner A, Amorim P, Fransoo JC, Honhon D, Kuhn H, Martinez de Albeniz V, Robb D (2021) Digitalization and omnichannel retailing: Innovative OR approaches for retail operations. *European Journal of Operational Research* 294(3):817–819.
- KC D, Terwiesch C (2009) Impact of workload on service time and patient safety: An econometric analysis of hospital operations. *Management Science* 55(9):1486–1498.
- Kesavan S, Lambert SJ, Williams JC, Pendem PK (2022) Doing well by doing good: Improving retail store performance with responsible scheduling practices at the Gap, Inc. *Management Science* 68(11):7818–7836.
- Kremer M (1993) The O-ring theory of economic development. *Quarterly Journal of Economics* 108(3):551–575.
- Krusell P, Ohanian LE, Ríos-Rull JV, Violante GL (2000) Capital-skill complementarity and inequality: A macroeconomic analysis. *Econometrica* 68(5):1029–1053.
- Martinez-Moyano IJ, McCaffrey DP, Oliva R (2014) Drift and adjustment in organizational rule compliance. *Organization Science* 25(1):44–65.
- Mas A, Moretti E (2009) Peers at work. *American Economic Review* 99(1):112–145.
- McKnight DH, Choudhury V, Kacmar C (2002) Developing and validating trust measures for e-commerce: An integrative typology. *Information Systems Research* 13(3):334–359.
- Milgrom P, Roberts J (1990) The economics of modern manufacturing: Technology, strategy, and organization. *American Economic Review* 80(3):511–528.
- Mittal V, Ross WT, Baldasare PM (1998) The asymmetric impact of negative and positive attribute-level performance on overall satisfaction and repurchase intentions. *Journal of Marketing* 62(1):33–47.
- Oliva R, Sterman JD (2001) Cutting corners and working overtime: Quality erosion in the service industry. *Management Science* 47(7):894–914.
- Oliva R, Sterman JD, Giese M (2003) Limits to growth in the new economy: Exploring the ‘get big fast’ strategy in e-commerce. *System Dynamics Review* 19(2):83–117.
- Rahmandad H, Repenning N (2016) Capability erosion dynamics. *Strategic Management Journal* 37(4):649–672.
- Repenning NP, Sterman JD (2001) Nobody ever gets credit for fixing problems that never happened. *California Management Review* 43(4):64–88.
- Saghiri S, Wilding R, Mena C, Bourlakis M (2017) Toward a three-dimensional framework for omni-channel. *Journal of Business Research* 77:53–67.

- Sirdeshmukh D, Singh J, Sabol B (2002) Consumer trust, value, and loyalty in relational exchanges. *Journal of Marketing* 66(1):15–37.
- Sterman JD (2000) *Business Dynamics: Systems Thinking and Modeling for a Complex World*. Boston: Irwin/McGraw-Hill.
- Ton Z (2014) *The Good Jobs Strategy: How the Smartest Companies Invest in Employees to Lower Costs and Boost Profits*. New York: Houghton Mifflin Harcourt.
- Tucker AL (2007) An empirical study of system improvement by frontline employees in hospital units. *Manufacturing & Service Operations Management* 9(4):492–505.
- Webb MD (2023) Reworking wild bootstrap-based inference for clustered errors. *Canadian Journal of Economics* 56(3):839–858.